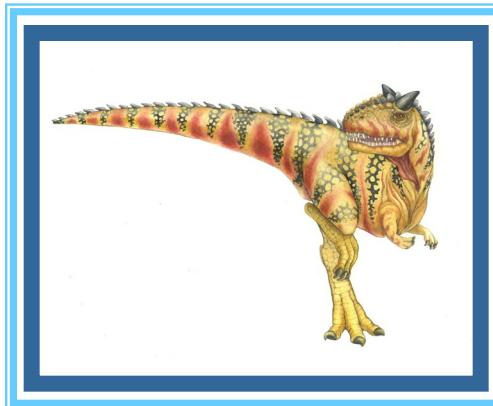


Chapter 3: Processes





Process Concept

- Early computers allowed only **one program to be executed at a time.**
 - This program had complete control of the system and access to all its **resources**.
 - Contemporary computer systems allow **multiple programs** to be **loaded into memory** and executed concurrently on a single CPU system.
 - This evolution resulted in the notion of a **process**
 - which is a program in execution (means *executable file of the program in the main memory, the CPU scheduling algorithm controls that*).
 - *A process is the unit of work in a modern computing system.*
-





Process Concept

- Therefore, a system consists of a collection of **processes**, some *executing user code*, others *executing operating system code*.
- Potentially, all these **processes** can execute concurrently, with a single CPU (or multiple CPUs).
- An OS executes a variety of programs:
 - **Batch system** – jobs (assigned work)
 - **Time-shared systems** – user programs or tasks
 - Even on a *single-user system*, a user may be able to run several programs at one time:
 - ▶ a word processor, a Web browser, and an e-mail package.





The Process

- **Process** – a program in execution (its .exe file in main memory, and it is taken by the CPU for execution or waiting for execution); process execution must progress sequentially.
- A **process** has multiple parts:
 - The **text section** is the **.exe file**
 - **Stack** stores *temporary data when invoking functions*
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time
 - ▶ The structure of a process in memory is shown in **Figure 3.1.**





The Process




Figure 3.1 Process in memory.





The Process

- Notice that the sizes of the **text** (.exe file section) and **data** sections of the memory are **fixed**, as their sizes do not change during **program run time**.
- The **stack** and **heap** sections can shrink and grow dynamically during **program execution (run time)**.
 - The **heap** will grow as memory is *dynamically allocated*, and will shrink when memory is *returned to the system*.
- Each time a **function is called**, an **activation record** containing *function parameters*, *local variables*, and the *return address* is pushed onto the **stack**;
 - When control is returned from the function, the activation record is popped from the stack.





The Process

- A processor register that supports program execution is the **program counter** register.
- It is emphasized that a program by itself is not a process.
- A **program** is a **passive entity**, such as a file containing a list of instructions stored on **disk** (often called an **.exe file**).
- A **process** is an **active entity**, with a **program counter** specifying the next instruction to execute and a set of associated resources.
 - A **program becomes a process** when an **executable file** is loaded into memory.
 - Two common techniques for **loading executable files** are double-clicking an **.exe icon** of the program on a disk drive or entering the name of the **executable file** on the command line.





The Process

- Although **two processes** may be associated with the same program, they are considered two separate execution sequences.
 - For instance, several users may be running different copies of the mail program, or the same user invokes many copies of the web browser, which are **separate processes**.
- A process can itself be an execution environment for other code.
 - The **Java programming environment** provides a good example.
 - In most cases, a Java program runs within the **Java Virtual Machine (JVM)**.
 - The **JVM** executes as a process that interprets the **loaded Java code** and takes actions (via native machine instructions) on its behalf.
 - For example, to run the compiled Java program **Program.class**, we would enter: **java Program**





Process State

- As a **process** executes, it changes state:
 - The state of a process is defined in part by the current activity of that process.
- A **process** may be in one of the following **states**:
 - **New**: The process is being created.
 - **Running**: Instructions are being executed.
 - **Waiting**: The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
 - **Ready**: The process is waiting to be assigned to a processor.
 - **Terminated**: The process has finished execution.





Process State

- It is important to realize that only **one process can be *running* on any processor core at any instant**. Many processes may be ***ready*** and ***waiting***, however.
- The state diagram corresponding to these states is presented in **Figure 3.2**.




Figure 3.2 Diagram of process state.





Process Control Block

- Each process is represented in the OS by a **process control block (PCB)**.
- A **PCB** is shown in **Figure 3.3**.
- In brief, the **PCB** serves as the repository for all the data needed to start, or restart, a process, along with some accounting data.




Figure 3.3 Process control block (PCB).





Process Control Block

- A **PCB** contains many pieces of information associated with a specific **process**:
 - **Process state:** The state may be new, ready, running, waiting, halted, and so on.
 - **Program counter:** The counter indicates the address of the next instruction to be executed for this process.
 - **CPU registers:** The size of registers depends on the computer architecture.
 - **CPU-scheduling information:** This information includes a process priority, pointers to scheduling queues, etc.
 - **Memory-management information:** This information includes the value of the *base* and *limit* registers and the *page tables*, etc
 - **Accounting information:** This information includes the amount of CPU and real time used, time limits, account numbers, job or process numbers, etc.
 - **I/O status information:** This information includes the list of I/O devices allocated to the process, a list of open files, and so on.





Process Scheduling

- The **objective of multiprogramming** is to have some process always running to **maximize CPU utilization**.
- The **objective of time sharing (Round Robin method)** is to **switch** a CPU core among processes so frequently that users can interact with each program while it is running (see **Figure 3.4**).
 - To meet these objectives, the **CPU scheduler** selects a set of several available processes from the **Ready Queue** for program execution on **a core**.
 - For a system with a **single CPU core**, there will be **one process** running at a time.
 - *The number of processes currently in memory is known as the degree of multiprogramming.*





A CPU Switches from Process to Process




Figure 3.4 Diagram showing CPU switch from process to process.





Threads

- The process model discussed so far has implied that a process is a program that performs a single **thread** of execution.
- This **single-threaded control** allows the process to perform only one task at a time.
- Most modern operating systems have extended the **process concept to multiple threads of execution**, enabling it to perform more than one task at a time.
- This feature is beneficial only on **multicore systems**, where multiple threads can run in parallel.
- A **multithreaded word processor** could assign one thread to manage user input while another thread runs the spell checker.
- On systems that support **thread execution**, the **PCB** is expanded to include information for each thread called a **TCB** (thread control block).
 - Other changes throughout the system are also needed to support threads. **Chapter 4 explores threads in detail.**





Scheduling Queues

- As processes are created, they are put into a **ready queue**, where they are **ready and waiting to execute on a CPU's core**.
 - This queue is generally stored as a **linked list**; a **ready-queue header contains pointers to the first PCB in the list**, and each **PCB includes a pointer field that points to the next PCB in the ready queue** (**Figure 3.4**).
 - When a process is allocated a **CPU core**, it executes for a while. Eventually, it terminates, is interrupted, or waits for a particular event, such as the completion of an **I/O request**.
 - Suppose the process makes an **I/O request** to a device such as a disk.
 - Since **I/O devices** run significantly slower than processors, the process will have to **wait for the I/O** to become available. Processes that are waiting for a specific event — such as **I/O completion** — are placed in a **wait queue** (**Figure 3.4**).





Scheduling Queues

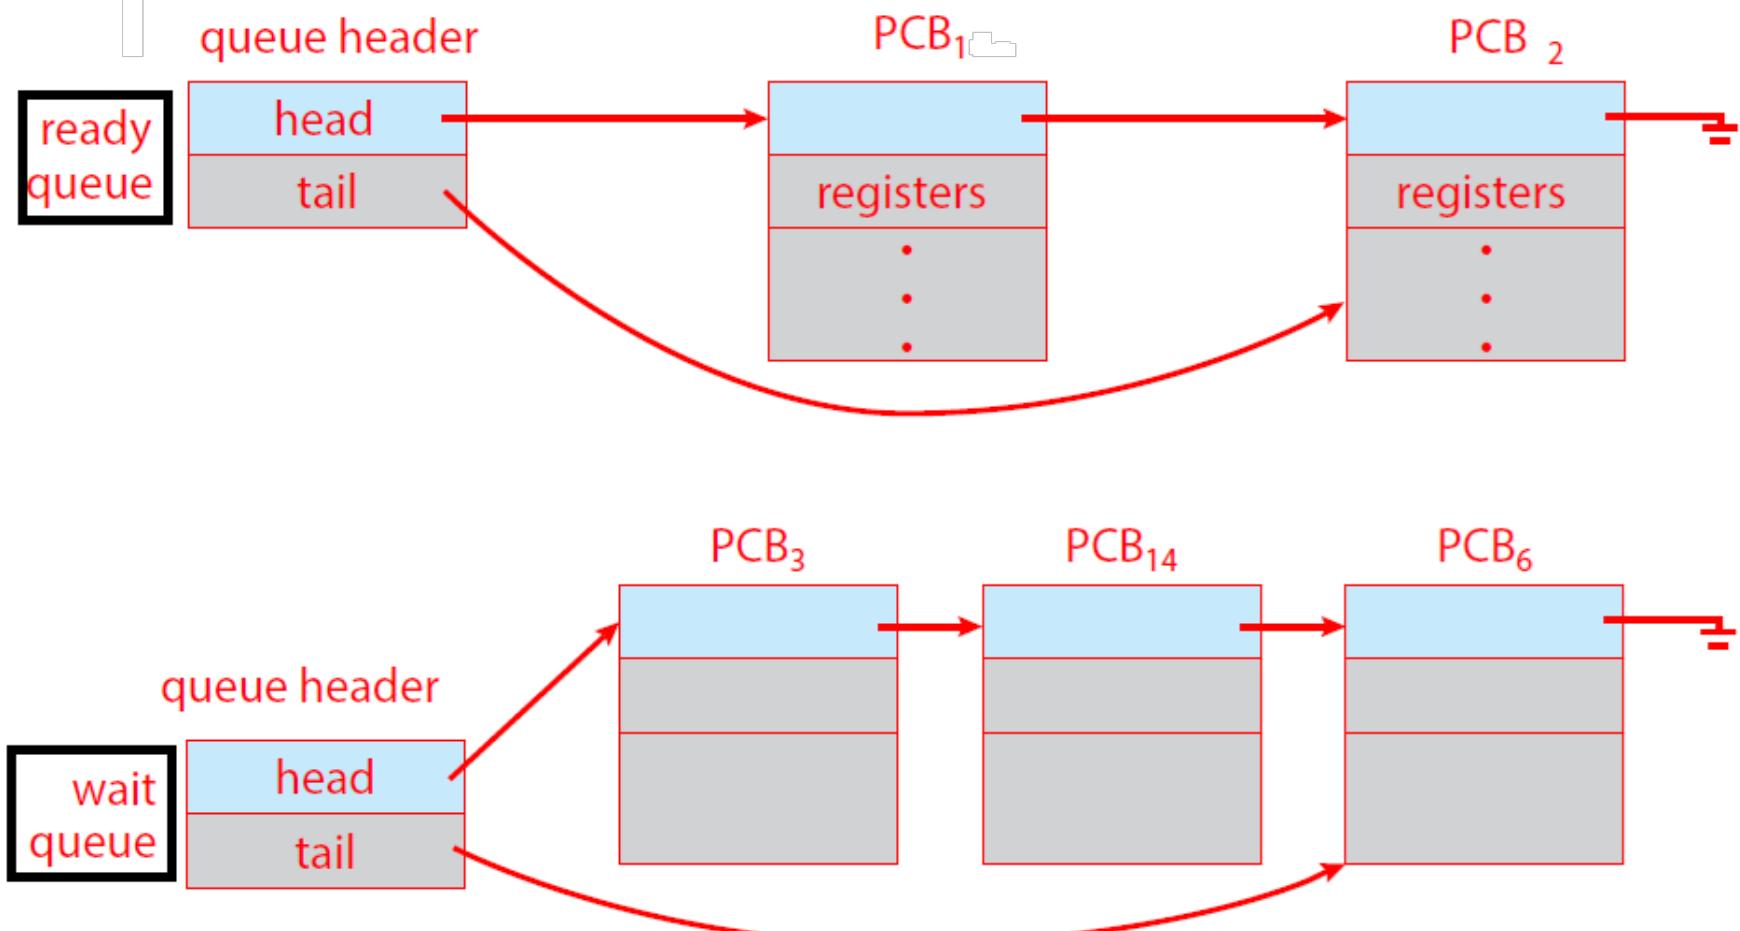


Figure 3.4 The ready queue and wait queues.





Scheduling Queues

- A common representation of process scheduling is a **queueing diagram** (in **Figure 3.5**).
- **Two types of queues** are present: the **ready queue** and a set of **wait queues**.
 - The **circles** represent the *resources* that serve the queues, and the **arrows** indicate *the flow of processes* in the system.
 - A **new process** is initially put in the **ready queue**. It waits there until it is selected for **execution** or **dispatched**.





Scheduling Queues

- Once the **process** is allocated a **CPU core**, one of several events could occur:
 - The process could issue an **I/O request** and then be placed in an **I/O wait queue**. The process could create a **new child process** and then be placed in a **wait queue**.
 - The process could be removed from the core, because of an **interrupt** or its **time slice expires**, and be put back in the **ready queue**.
- In the first two cases (in **Figure 3.5**), the process eventually switches from the **waiting state** to the **ready state** and is then put back in the **ready queue**.
 - A process continues this cycle until it **terminates**.
 - The **PCB and resources** of the terminated process are **deallocated**.





Scheduling Queues




Figure 3.5 Queueing-diagram representation of process scheduling.





CPU Scheduling

- A **process** migrates among the **ready queue** and **various wait queues** throughout its lifetime.
- The role of the CPU scheduler is to select a process from the ready queue and allocate a CPU core to it.
- The **CPU scheduler** must frequently select a **new process** for the CPU.
- An **I/O-bound process** may execute for only a **few milliseconds** before waiting for an **I/O request**.
- Although a **CPU-bound process** will require a **CPU core**, it forcibly removes it from the **CPU** and places it in the **ready queue**, scheduling another process to run (due to the **preemptive nature** of the scheduler).





CPU Scheduling

- Therefore, the **CPU scheduler executes** at least once every **100 milliseconds**, although typically much more frequently.
 - A process may execute for only a **few milliseconds** before waiting for an **I/O request**.
 - Often, the **preemptive scheduler** executes processes at least once every 100 milliseconds.
 - If it takes **10 milliseconds** to decide to execute a process for **100 milliseconds** (total is 110ms), then $10/(110) = 9\%$ of the CPU is being used (wasted) simply for scheduling the work.





I/O-bound process and CPU-bound process

- In general, most processes can be described as either **I/O bound** or **CPU bound**.
 - An **I/O-bound process** is one that spends most of its time doing I/O than it spends doing computations.
 - A **CPU-bound process**, generates less I/O requests and using more of its time doing computations.

```
{  
    printf("\nEnter the first integer: ");  
    scanf("%d", &a);  
    printf("\nEnter the second integer: ");  
    scanf("%d", &b);  
  
    c = a+b  
    d = (a*b)-c  
    e = a-b  
    f = d/e  
}  
  
printf("\n a+b= %d", c);  
printf("\n (a*b)-c = %d", d);  
printf("\n a-b = %d", e);  
printf("\n d/e = %d", f);  
}
```

} I/O cycle

} CPU cycle

} I/O cycle





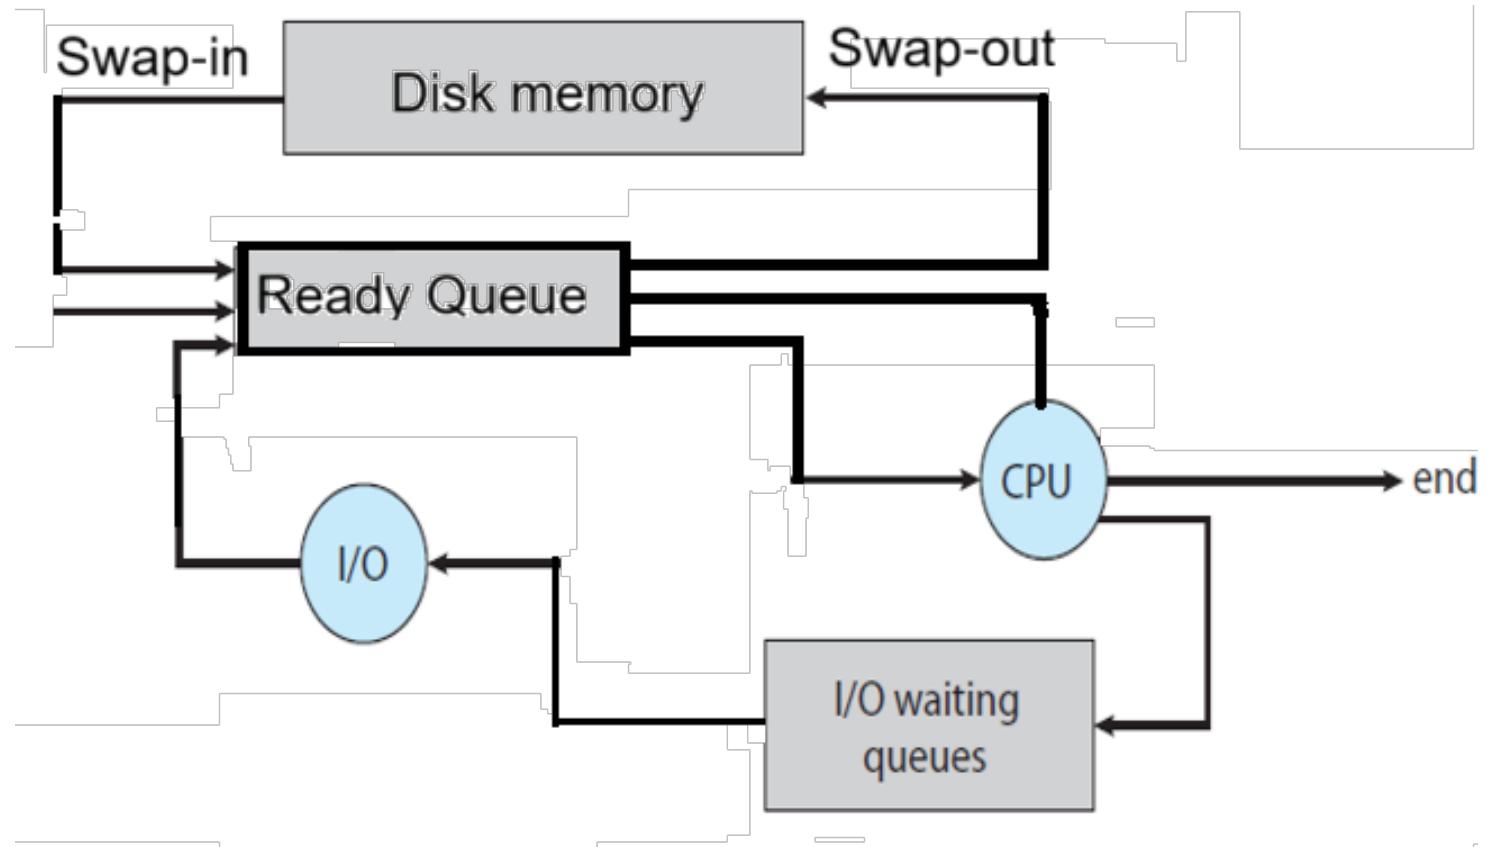
Process Swapping

- Some operating systems have an **intermediate form of scheduling**, known as **swapping**
 - The key idea of **swapping** is to remove a **process** from **memory (ready queue)** to **disk memory (virtual memory)** and thus *reduce the degree of multiprogramming*.
 - Later, the process can be reintroduced into memory, and its execution can be continued where it left off.
 - This scheme is known as **swapping** because a process can be “**swapped out**” from **memory to disk**, where its status is saved, and later “**swapped in**” from **disk back to memory**, where its status is restored.
 - *Swapping is typically only necessary when memory has been overcommitted and must be freed up.*





Process Swapping





Context Switch

- **Interrupts** cause the operating system to change a **CPU core** from its current task and to run a **kernel routine**.
 - *Such operations happen frequently on general-purpose systems.*
 - When **an interrupt occurs**, the system needs to save the current **context** of the process running on the CPU core so that it can restore that context when its processing is done, essentially suspending the process and then resuming it.
- The **context of a process** is represented in the form of a **PCB** of the OS.
 - PCB includes the value of the **CPU registers**, the **process state** (see Figure 3.2), and **memory-management information**.
 - A **state save** of the **current state of the CPU core**, be it in **kernel or user mode**, and then a **state restore** to resume operations.





Context Switch

- **Switching the CPU core to another process** requires performing a state save of the current process and a state restore of a different process.
- This task is known as a **context switch** and is illustrated in **Figure 3.6**.
- **When a context switch occurs;** *the kernel saves the context of the old process in its PCB and loads the saved context of the new process scheduled to run.*
- **Context switch time is pure overhead**, because the system does no useful work while switching.
- **Switching speed** varies from machine to machine, depending on the ***memory speed, the number of registers*** that must be copied, and the existence of special instructions (such as ***load*** or ***store*** instructions).
- Context-switch times are highly dependent on **hardware support**.






Figure 3.6 Diagram showing context switch from process to process.





Process Creation (1)

- During the course of execution, a **process** may create several new processes (called **spawning**).
- The creating process is called a **parent process**, and the new processes are called the **children processes**.
 - Forming a **tree** of processes.
 - Most operating systems including **UNIX**, **Linux**, and **Windows** identify processes according to a unique **process identifier** (or **pid**),
- The **pid** provides a unique value for each process in the system,
 - used as an index to access various attributes of a process within the kernel.





Process Creation (2)

- Figure 3.8 illustrates a typical process tree for the Linux operating system, showing the name of each process and its pid.

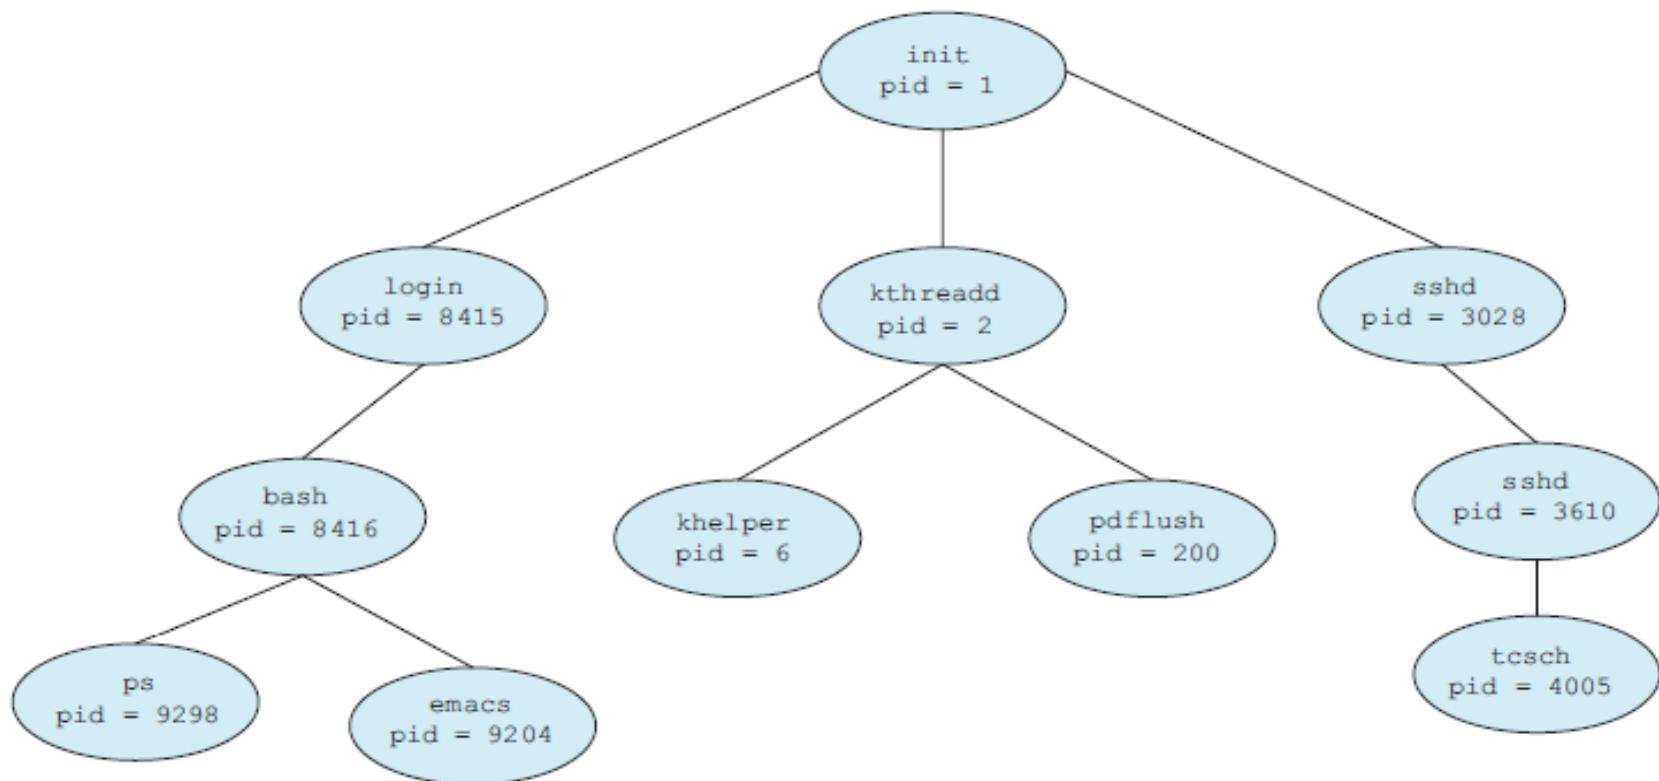


Figure 3.8 A tree of processes on a typical Linux system.





Process Creation (3)

- The **init** process (which always has a pid of 1) serves as the root parent process for all user processes.
- Once the system has booted, the **init** process can also create various user processes, such as a web or print server, an **ssh** server, and the like.
- The **kthreadd** process is responsible for creating additional processes that perform tasks on behalf of the kernel (in this situation, **khelper** and **pdflush**).
- The **sshd** process is responsible for managing clients that connect to the system by using **ssh** (which is short for **secure shell**).
- The **login** process is responsible for managing clients that directly log onto the system.
- In this example, a client has logged on and is using the **bash** shell, which has been assigned pid 8416.
- Using the **bash** command-line interface, this user has created the process **ps** as well as the **emacs** editor.
- On UNIX and Linux systems, we can obtain a listing of processes by using the **ps** command. For example, the command **ps -el**





Process Termination

- Process executes last statement and then asks the OS to delete it using the exit() system call.
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are de-allocated by operating system
- Parent may terminate the execution of children processes using the abort() system call. Some reasons for doing so:
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates.
 - A process can cause the termination of another process via an appropriate system call (for example **TerminateProcess()** in **Windows**).

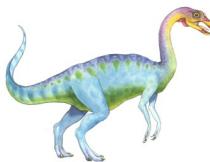




Independent and Cooperating processes

- Processes executing concurrently in the OS may be either ***independent*** processes or ***cooperating processes***.
 - A process is ***independent*** if it cannot affect or be affected by the other processes executing in the system.
 - ▶ Any process that does not share data with any other process is independent.
 - A process is ***cooperating*** if it can affect or be affected by the other processes executing in the system.
 - ▶ Clearly, any process that shares data with other processes is a ***cooperating process***.






Process Cooperation

- There are several reasons for providing an environment that allows **process cooperation**:
 - **Information sharing:** several users may be interested in the same piece of information (for instance, a shared file), system must allow concurrent access to such information.
 - **Computation speedup:** If we want a particular task to run faster, we must break it into subtasks and allow them to execute in parallel.
 - **Modularity:** We may want to construct the system in a modular fashion, dividing the system functions into separate processes or threads.
 - **Convenience:** Even an individual user may work on many tasks at the same time. For instance, a user may be editing, listening to music, and compiling in parallel.





Inter-process Communication (IPC)

- Cooperating processes require an inter-process communication (IPC) mechanism that will allow them to exchange data and information.
- There are two fundamental models of IPC:
 - shared memory: In the shared-memory model, a region of memory that is shared by cooperating processes is established.
 - ▶ Processes can then exchange information by reading and writing data to the shared region.
 - Message passing: In the message-passing model, communication takes place by means of messages exchanged between the cooperating processes.
- The two communications models are contrasted in **Figure 3.12**.





Interprocess Communication




Figure 3.12 Communications models. (a) Message passing. (b) Shared memory.





Interprocess Communication

- **Message passing** is useful for exchanging smaller amounts of data, because no conflicts need be avoided.
 - It is easier to implement in a **distributed system** than shared memory.
 - It needs system calls to implement .
 - Hence it requires more time-consuming task of kernel intervention.
- **Shared memory** can be faster than message passing, since its does not need system calls to implement.
 - Here system calls are required only to establish shared memory regions.
 - ▶ Once shared memory is established, all accesses are treated as routine memory accesses, and no assistance from the kernel is required.

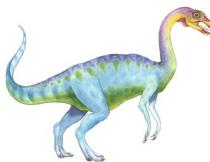




Shared-Memory Systems

- The OS tries to prevent one process from accessing another process's memory.
 - **Shared memory** requires that two or more processes agree to remove this restriction.
 - They can then exchange information by reading and writing data in the shared areas.
 - The form of the data and the location are determined by these processes and are not under the operating system's control.
 - The processes are also responsible for ensuring that they are not writing to the same location simultaneously.





Shared-Memory Systems

- To illustrate the concept of **cooperating processes**, let's consider the **producer–consumer** problem, which is a common paradigm for cooperating processes.
 - A **producer process** produces information that is consumed by a consumer process.
 - ▶ For example, a **compiler** may produce **assembly code** that is consumed by an **assembler** section of the compiler.
 - ▶ The assembler, in turn, may produce **object modules** that are consumed by the **loader** section of the assembler.
- The **producer–consumer problem** also provides a useful metaphor for the **client–server** paradigm.

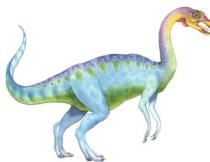




producer–consumer problem

- One solution to the **producer–consumer** problem uses **shared memory**.
 - To allow producer and consumer processes to run concurrently, we must have available **a buffer** of items that can be filled by the producer and emptied by the consumer.
 - This **buffer** will reside in a region of memory that is **shared** by the **producer** and **consumer** processes.
 - A **producer** can produce one item while the **consumer** is consuming another item.
 - The **producer** and **consumer** must be synchronized, so that the consumer does not try to consume an item that has not yet been produced.





producer–consumer problem

- Two types of buffers can be used:
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size
- In this case, the consumer must wait if the buffer is empty, and the producer must wait if the buffer is full.
- The code for the producer process is shown in **Figure 3.13**, and the code for the consumer process is shown in **Figure 3.14**.





producer–consumer problem

```
item next_produced;

while (true) {
    /* produce an item in next_produced */

    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Figure 3.13 The producer process using shared memory.





producer–consumer problem

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```

Figure 3.14 The consumer process using shared memory.





Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send(message)**
 - **receive(message)**
- The *message size* is either fixed or variable





Message Passing (Cont.)

- If processes **P** and **Q** wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via **send/receive**
- **Implementation issues:**
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?





Message Passing (Cont.)

□ Implementation of communication link

□ **Physical:**

- Shared memory
- Hardware bus
- Network

□ **Logical:**

- Direct or indirect
- Synchronous or asynchronous
- Automatic or explicit buffering





Direct Communication

- Processes must name each other explicitly:
 - **send (P, message)** – send a message to process P
 - **receive(Q, message)** – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication - Mailbox

- Messages are directed and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique **id**
 - Processes can communicate only if they share the common mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional





Indirect Communication - Mailbox

□ Operations

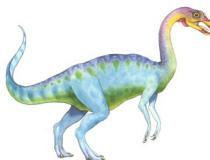
- **create** a new mailbox (port)
- **send** and **receive** messages through mailbox
- **destroy** a mailbox

□ Primitives are defined as:

send(A, message) – send a message to mailbox A

receive(A, message) – receive a message from mailbox A





Indirect Communication

□ Mailbox sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 , sends; P_2 and P_3 receive
- Who gets the message?

□ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver.
 - ▶ Sender is notified who the receiver was.





Synchronization

- ❑ Communication between processes takes place through calls to **send()** and **receive()** primitives.
- ❑ There are different design options for implementing each primitive. Message passing may be either **blocking** or **non-blocking**;
- ❑ **Blocking** is considered **synchronous**
 - ❑ **Blocking send** -- the sender is blocked until the message is received
 - ❑ **Blocking receive** -- the receiver is blocked until a message is available
- ❑ **Non-blocking** is considered **asynchronous**
 - ❑ **Non-blocking send** -- the sender sends the message and continues
 - ❑ **Non-blocking receive** -- the receiver receives:
 - ❑ A valid message, or Null message
- ❑ Different combinations possible
 - ❑ If both send and receive are blocking, we have a **rendezvous**



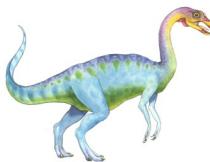


Synchronization (Cont.)

- With message passing, Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next_produced);  
    message next_consumed;  
} while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next consumed */  
}
```





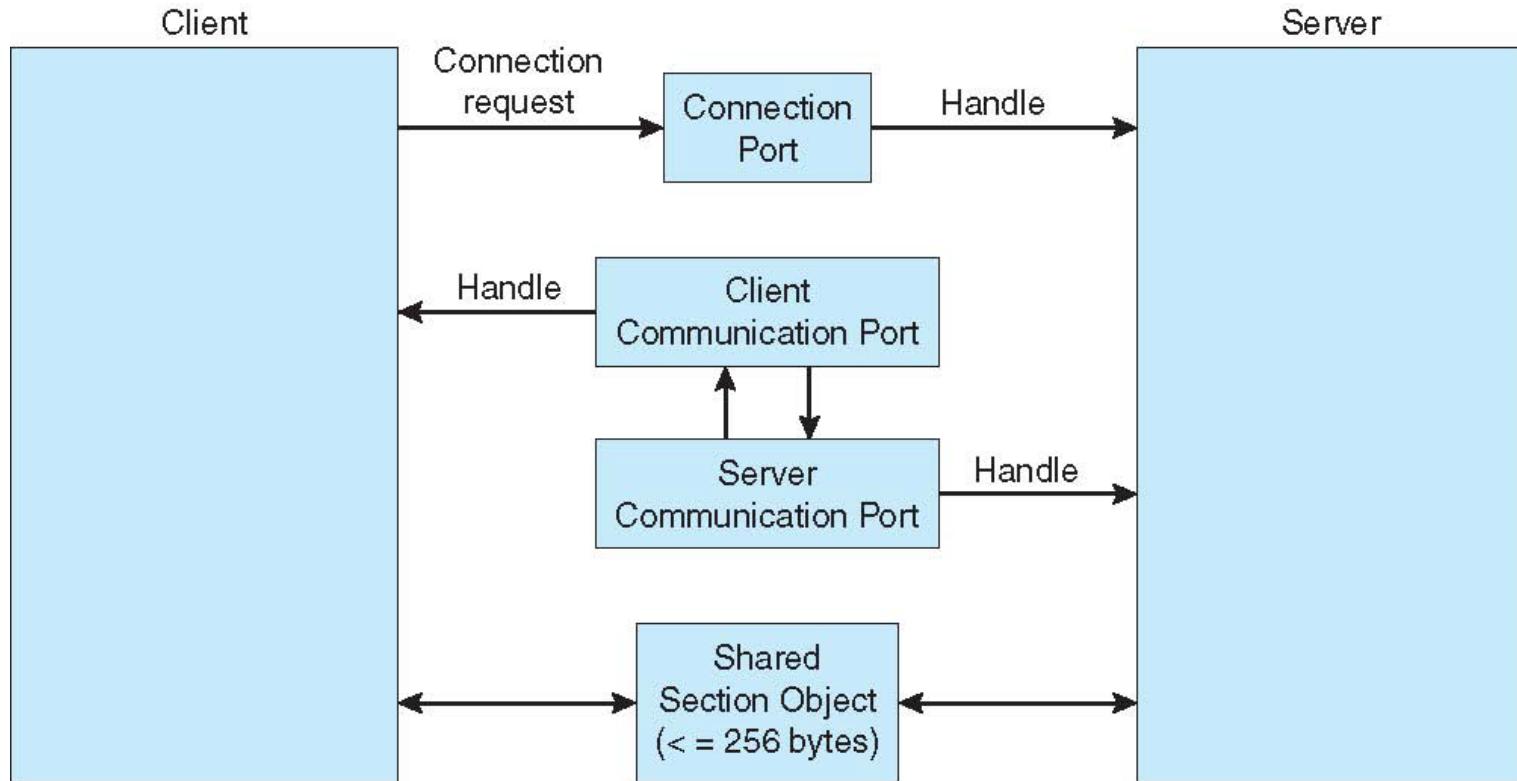
Examples of IPC Systems – Windows

- Message-passing centric via advanced **local procedure call (LPC)** facility
 - ***Only works between processes on the same system***
 - Uses ports (like **mailboxes**) to establish and maintain communication channels
 - Communication works as follows:
 - ▶ The client opens a handle to the subsystem's **connection port** object.
 - ▶ The client sends a **connection request**.
 - ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.





Local Procedure Calls in Windows





Remote Procedure Calls

- **Remote procedure call (RPC)** abstracts *procedure calls* between processes **on networked systems**
 - Again uses ports for service differentiation
- **Stubs** – are the intermediate data translation structures
- The **client-side stub** locates the server and **marshals** (bring together) the parameters
- The **server-side stub** receives this message, unpacks the marshaled parameters, and performs the procedure on the server
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**





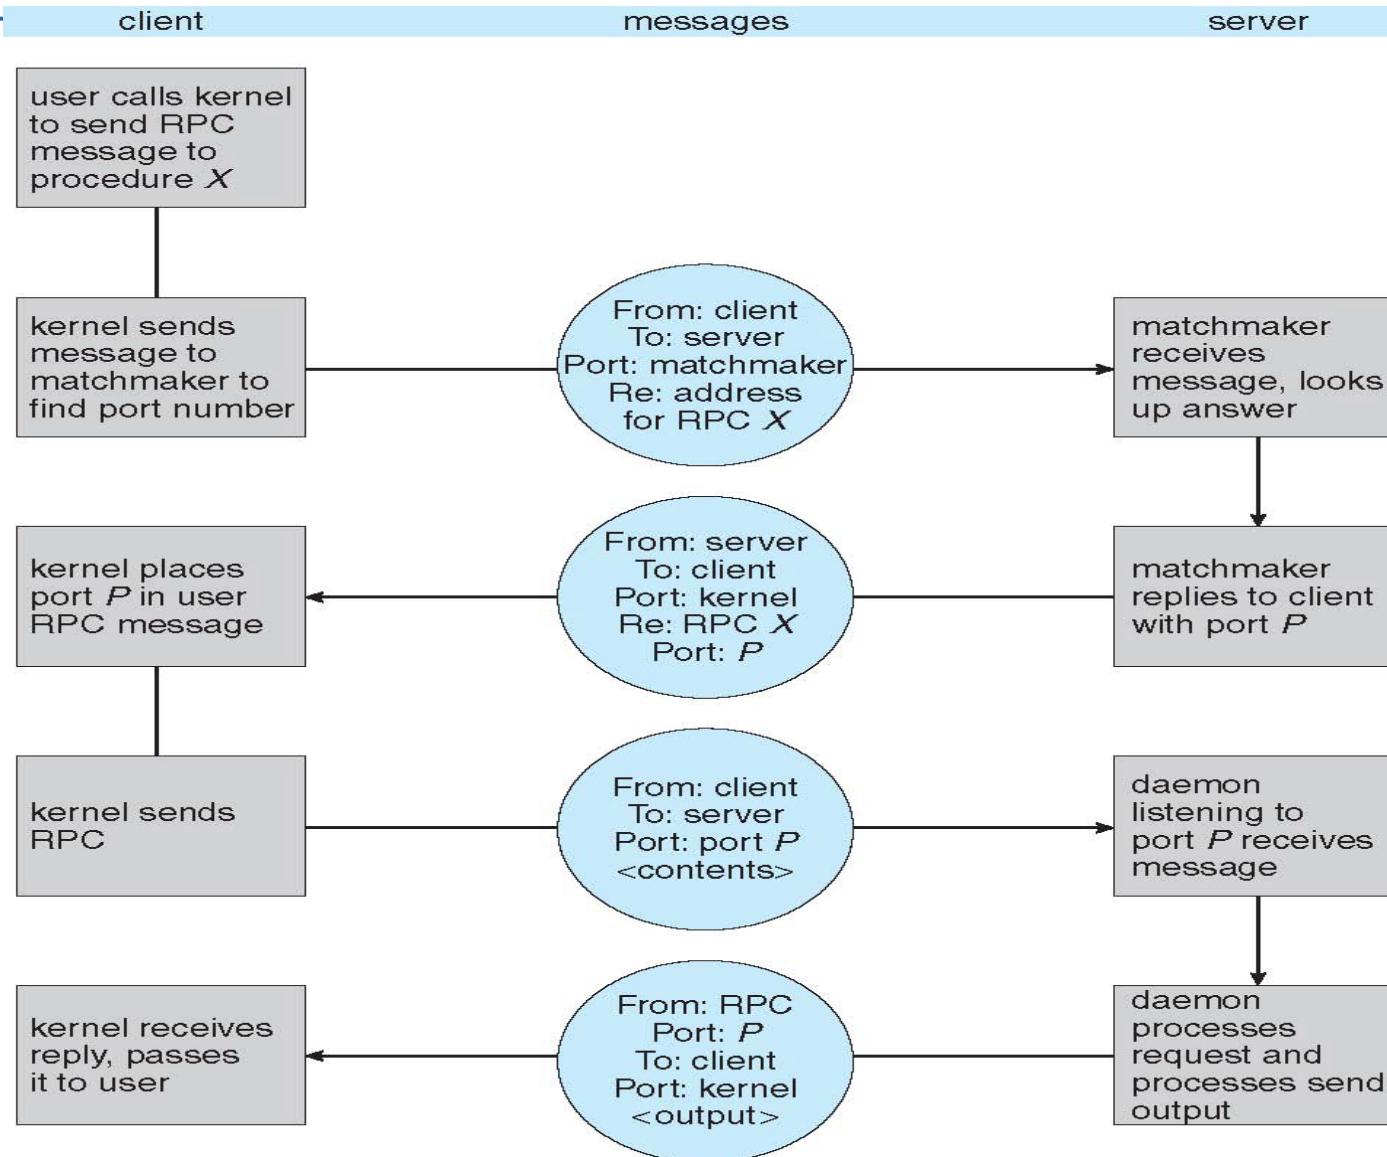
Remote Procedure Calls (Cont.)

- Data representation handled via **External Data Representation (XDL)** format to account for different architectures
 - **Big-endian** and **little-endian**
- Remote communication has more failure scenarios than local
 - Messages can be delivered ***exactly once*** rather than ***at most once***
- OS typically provides a **rendezvous** (or **matchmaker**) service to connect client and server





Execution of RPC





Programming Exercises

Show **thread creation** in Java using *Executor* and *Runnable* objects.

