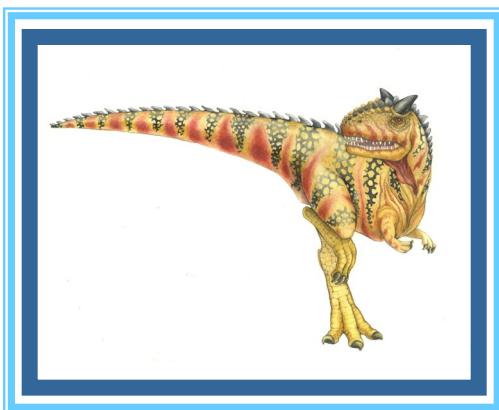


# Chapter 9: Memory Management

## Part II





# Fragmentation

---

- Both the **first-fit** and **best-fit** dynamic memory allocation strategies suffer from external fragmentation.
- External fragmentation exists when there is enough total memory space to satisfy an allocation request, but the available spaces are not contiguous and cannot be used:
  - As processes are loaded and removed from memory, the free memory space (hole) is broken into further little pieces.
  - The memory is fragmented into a large number of small unused holes.
- This external fragmentation problem can be severe.
  - If all these small pieces of memory holes were combined into a big free hole, it might be helpful for process allocation.





# Fragmentation

- There will be an issue of **internal fragmentation** as well:
  - Consider a **hole of size 18,464 bytes**. Suppose that the next process allocation request has a size of **18,462 bytes**. If the allocation exactly used the hole, there will be a **hole left with 2 bytes of size** —unused memory that is internal to a partition.
  - The *memory space allocated to a process is slightly larger than the process size, causing **internal fragmentation**.*
- The overhead to keep track of these unused holes resulting from **internal fragmentation** will be substantially larger than the hole itself.
  - **The general approach to avoiding internal fragmentation is to break the physical memory into fixed-sized blocks and allocate memory based on the block size.**





# Fragmentation

---

- One solution to the problem of external fragmentation is **compaction**:
  - *The compaction is to shuffle to place all free memory holes together in one large block.*
- **Compaction** is not always possible: If relocation is static and is done at assembly or load time, compaction cannot be done.
- The simplest **compaction** algorithm:
  - To move all processes toward one end of memory;
  - All holes move in the other direction, producing one large hole of available memory.
  - This scheme is expensive to implement.





# Fragmentation

---

- Another possible solution to the **external-fragmentation** is to permit the ***logical address space of processes to be noncontiguous***, thus allowing a process to be allocated physical memory wherever such memory is available.
- This is the strategy used in ***paging***, the most common memory-management technique for computer systems.
- Fragmentation is a general problem in computing that can occur wherever blocks of data are managed in memory.





# Paging

---

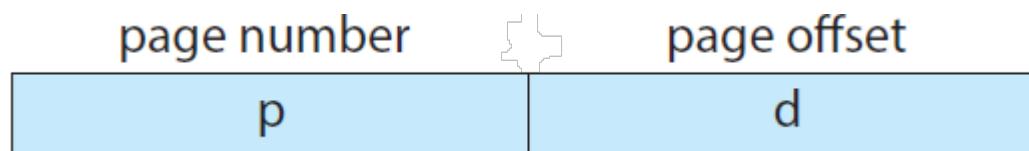
- **Paging** is a memory management scheme that permits a process's physical address space to be noncontiguous.
- **Paging** avoids **external fragmentation** and is implemented through cooperation between the operating system and the computer hardware.
- The basic method for implementing **paging** involves **breaking main memory into fixed-sized blocks** called frames and breaking **logical memory space** (disk memory space or program's memory) into **fixed-sized blocks** called pages.
- An **executing program** is converted into **pages** in its logical space and is **loaded into available memory frames** dynamically.





# Paging

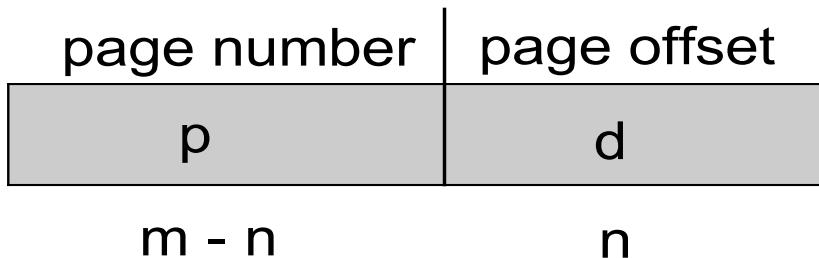
- For example, the **logical address space** (disk address space) is now totally separate from the **physical address space** (main memory address space), so a process can have a logical 64-bit address space even though the system has less than  $2^{64}$  bytes of physical memory space.
- In paging, every address generated by the CPU (called a **page address**) is divided into two parts:
  - a **page number** (**p**) and
  - a **page offset** (**d**):





# Paging

- The address generated by the CPU is divided into two:
  - **Page number ( $p$ )** – used as an **index** into a **page table**, which contains the **base address (frame address)** of each page in memory frame
  - **Page offset ( $d$ )** – combined with **base address** to define the **physical memory address** that is sent to the memory unit



If PA = 32bits ( $m$ ) and its offset size ( $n$ ) is 5 bits, then its page no. size is  $(m-n) = 27$  bits.

- For a given **logical address space**  $2^m$  and **page size**  $2^n$





# Paging

---

- The **page number (p)** is used as an index into a per-process **page table**. This is illustrated in **Figure 9.8**.
- The **page table** contains the **base address of each frame** in physical memory, and the **offset (d)** is the location in the frame being referenced.
- Thus, the **base address of the frame** is combined with the **page offset** to define the **physical memory address**.





# Paging Hardware

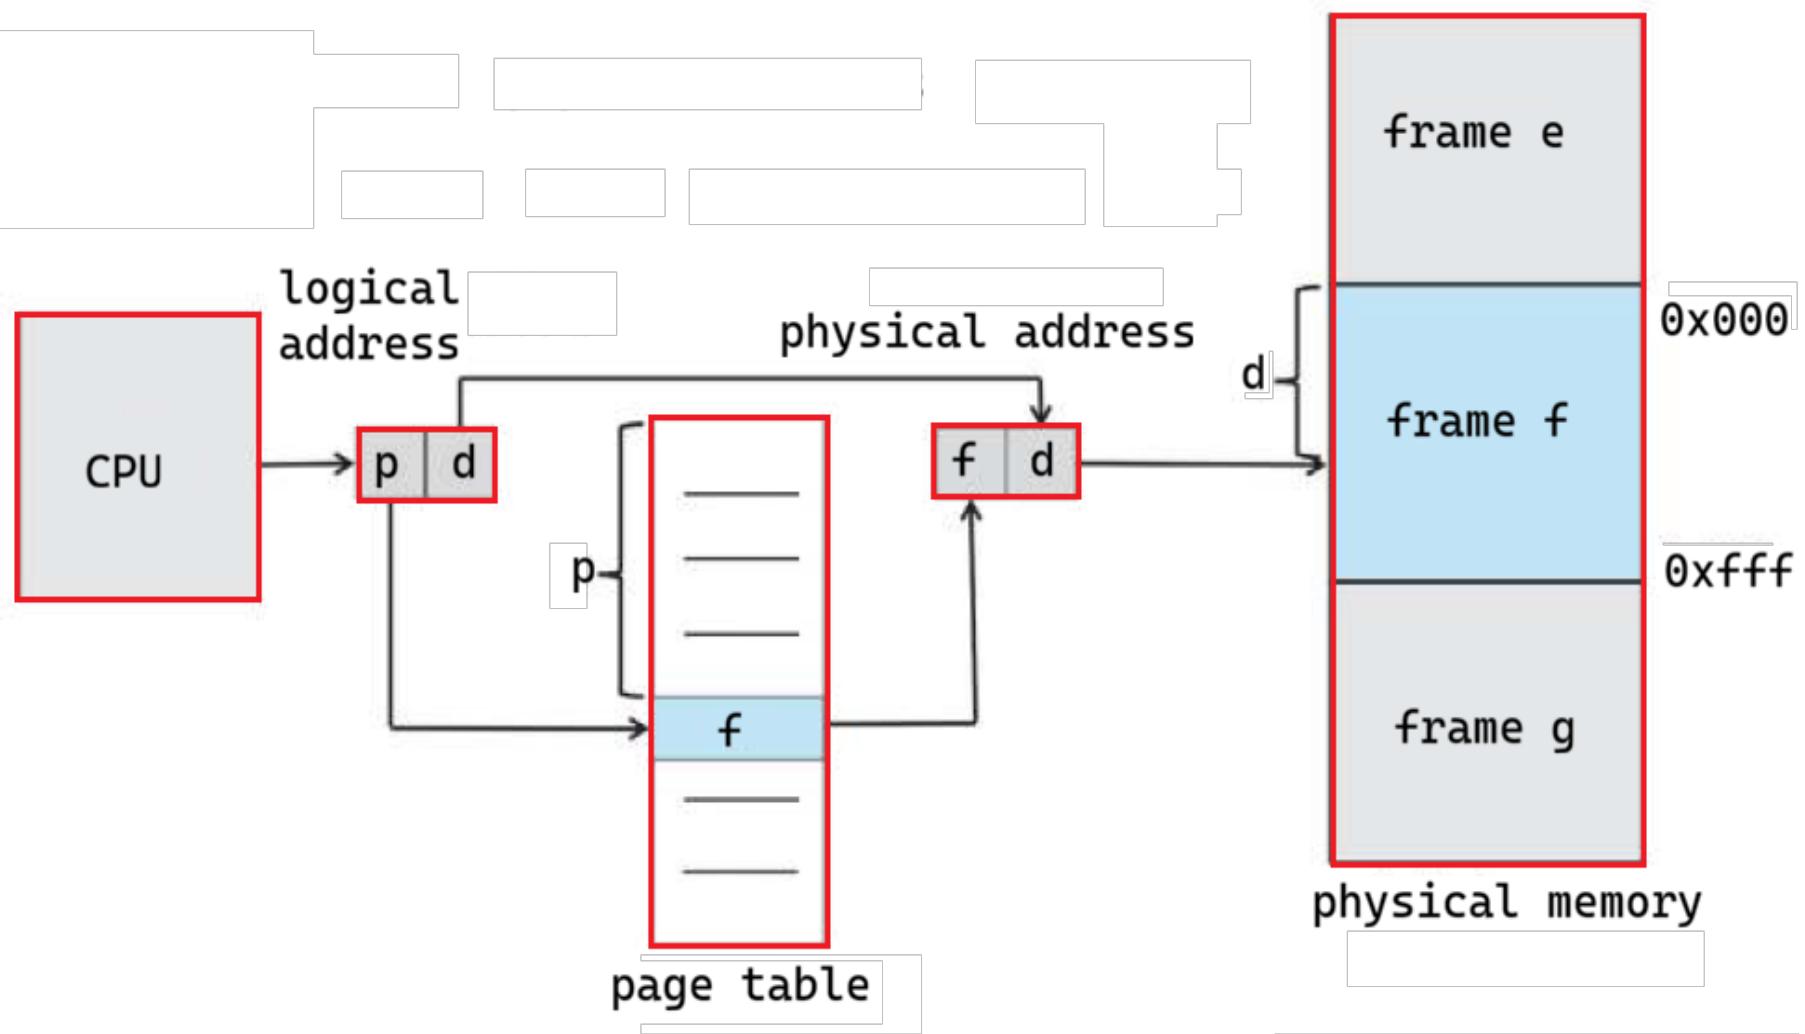


Figure 9.8 Paging hardware.





# Paging

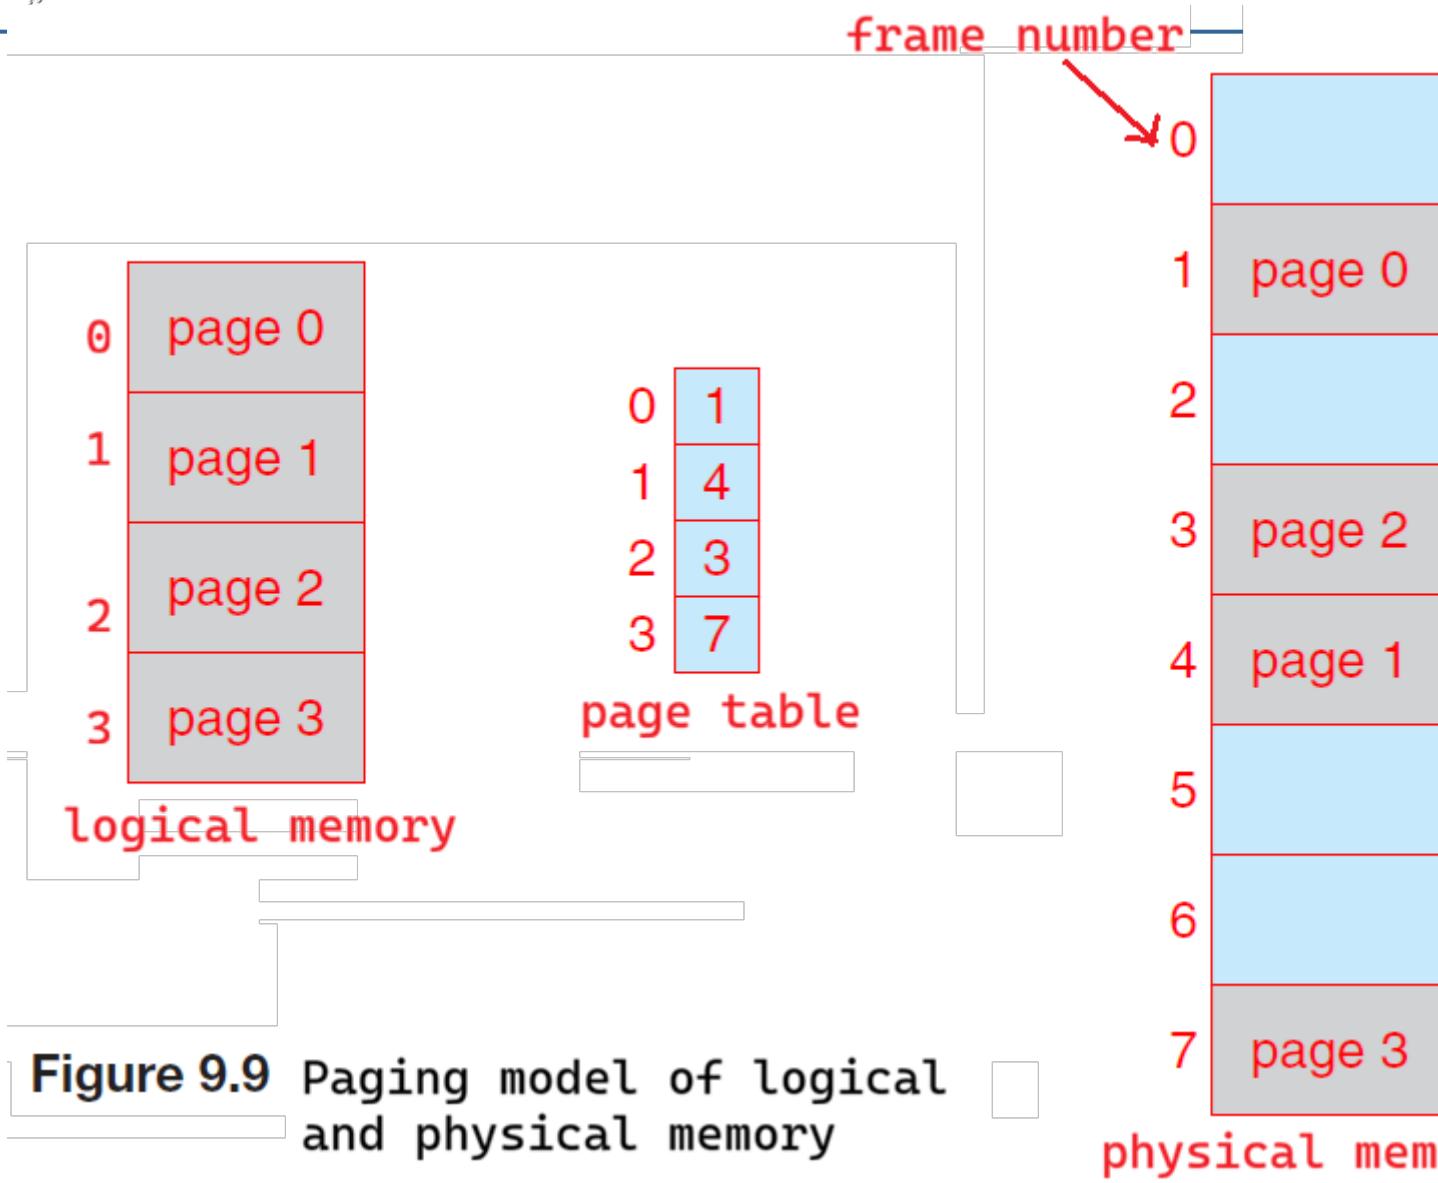
---

- The following outlines the steps taken by the MMU to translate a **logical address** generated by the CPU to a **physical address**:
  1. Extract the **page number  $p$**  and use it as an **index into the page table**.
  2. Extract the corresponding **frame number  $f$**  from the page table.
  3. Replace the **page number  $p$**  in the **logical address** with the **frame number  $f$** .
- As the **offset  $d$**  does not change, it is not replaced, and the **frame number** and **offset** now comprise the **physical address**.
- The paging model of memory is shown in **Figure 9.9**.





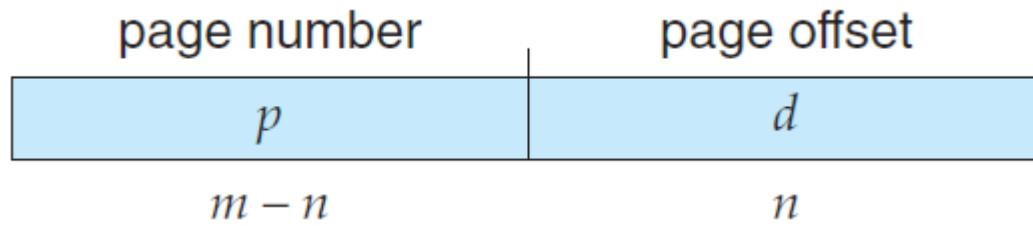
# Paging

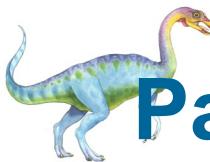




# Paging

- The **page size** (same as the **frame size**) is defined by the hardware.
  - The **size of a page is a power of 2**, typically varying between **4 KB** and **1 GB** per page, depending on the computer architecture.
- The selection of a **power of 2 as a page size** makes the **translation of a logical address into a page number and page offset** particularly easy.
  - If the **size of the logical address space is  $2^m$** , and a **page size is  $2^n$  bytes**, then the high-order of **( $m-n$ ) bits** of a logical address designate the **page number**, and the  **$n$  low-order bits** designate the **page offset**.
  - where  **$p$  is an index** into the page table, and  **$d$  is the displacement** within the page.

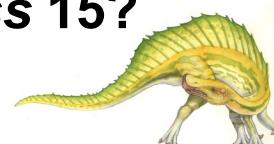




# Paging -Address Translation Scheme

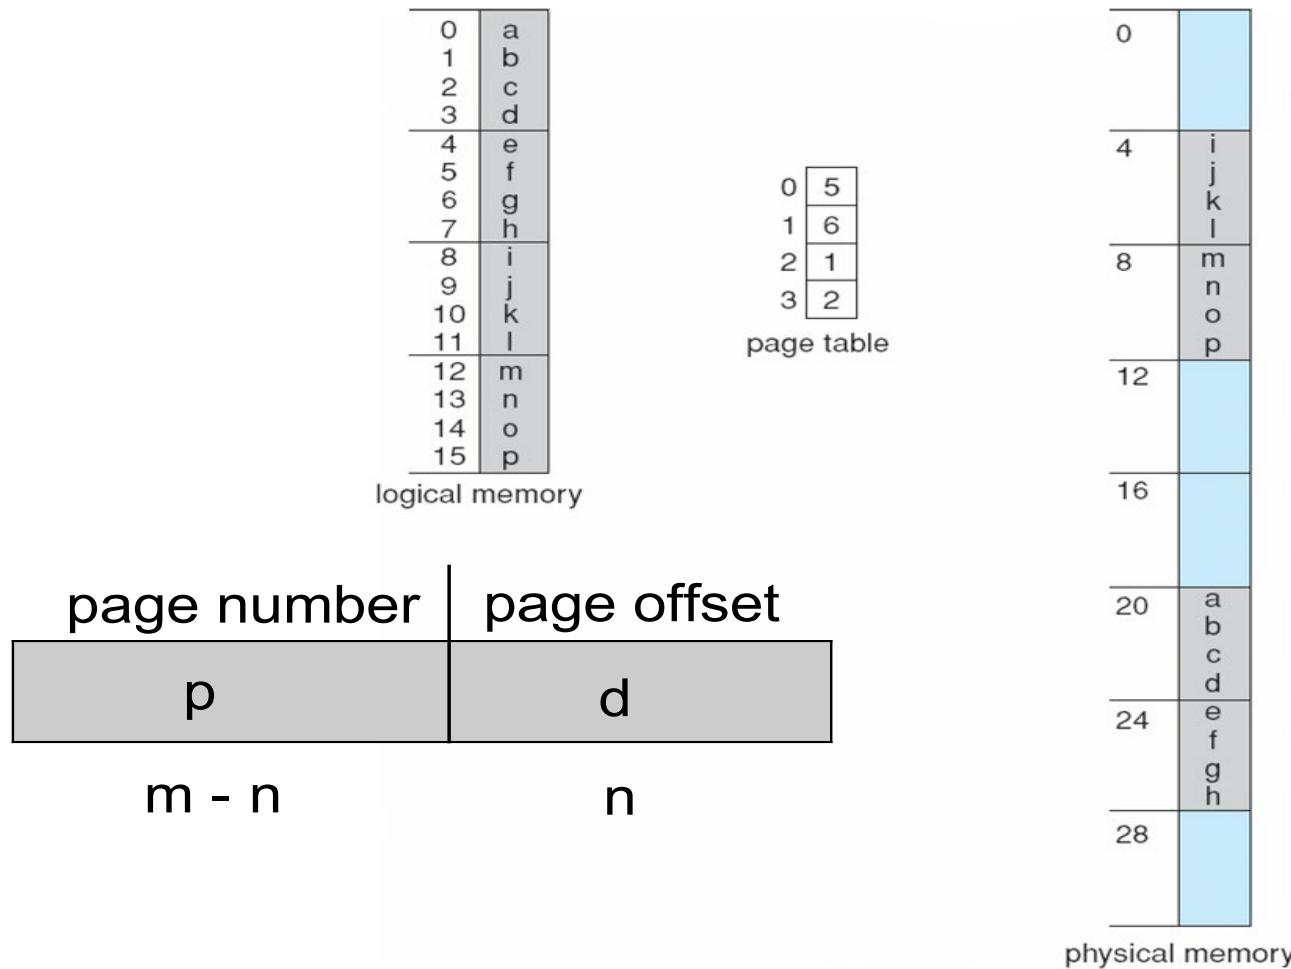
[Physical address = (frame no. × logical address-bits) + offset]

- Consider the memory in **Figure 9.10**. Here,  $n = 2$  and  $m = 4$ . Using a **page size of 4 bytes** and a physical memory of **32 bytes** ( $32 \times 8$ ) yields **8 pages**. We show how the programmer's view of memory can be mapped into physical memory.
- **Logical address 0** is (2-bit **page0**, 2-bit **offset 0 = 0000**). Indexing into the **page table**, notice that **page 0 is in frame 5**. Thus, **logical address 0** maps to **physical address 20** [=  $(5 \times 4) + 0$ ].
- **Logical address 3** (is 2-bit **page0**, 2-bit **offset 3 = 0011**) maps to physical address **23** [=  $(5 \times 4) + 3$ ].
- **Logical address 4** is (2-bit **page1**, 2-bit **offset 0 = 0100**); according to the page table, **page1** is mapped to **frame 6**. **Logical address 4** maps to physical address **24** [=  $(6 \times 4) + 0$ ].
- **Logical address 13** is (2-bit **page3**, 2-bit **offset 1 = 1101**) maps to physical address **9** [=  $(2 \times 4) + 1$ . *What about logical address 15?*





# Paging



$n = 2$  and  $m = 4$  32-byte memory and 4-byte pages

**Figure 9.10** Paging example for a **32-byte** memory with **4-byte** pages.





# Paging

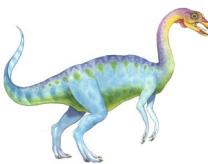
Consider a virtual memory system with a page size of **128-bit** and a physical memory of **256 bytes**. Each logical address has a **4-bit offset value**. The **page table** is shown below:

Page no.	Frame address
0	11
1	6
2	10
⋮	⋮
14	7
15	2

Based on the **page table**, calculate the **physical addresses** (in decimal) of the following **logical addresses**:

- 13.1. [1 point]  $01_{16}$
- 13.2. [1 point]  $2B_{16}$
- 13.3. [1 point]  $1A_{16}$
- 13.4. [1 point]  $FC_{16}$





# Paging

---

- Frequently, on a **32-bit CPU**, each **page-table entry** is **4 bytes** long, but that size can vary as well.
- A **32-bit entry** can point to one of **2<sup>32</sup> physical page frames**. If the frame size is **4 KB** (2<sup>12</sup>), then a system with **4-byte** entries can address 2<sup>44</sup> bytes (or 16 TB) of physical memory.





# Paging

---

- Consider a computer with a 32-bit logical address (VA) and a **4 Kbits** page size. How many entries are required in a page table? If each page table entry requires **4 bytes**, what is the total size of the page table?
- Solution:

number of pages = Logical memory (VM)size/page size =  
 $2^{32}/2^{12} = 2^{20}$

As there is one page table entry for each page, there are  $2^{20}$  entries. If each entry is **4 bytes**, the page table requires **2<sup>22</sup> bytes or 4M bytes** of memory.  
(that is,  $2^{20} \times 4$  byte = 4M bytes).





# Paging

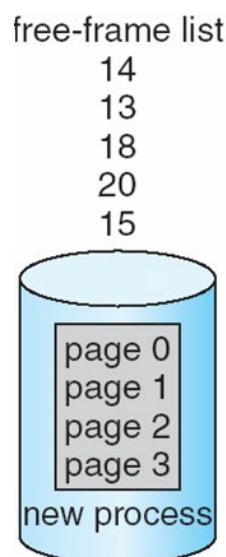
---

- When a process arrives in the system to be executed, its **size**, expressed in pages, is examined.
- Each **page** of the process needs one memory **frame**.
- Thus, if the process requires ***n* pages**, at least ***n* frames** must be available in memory.
- If ***n frames*** are available, they are allocated to this arriving process.
- The **first page** of the process is loaded into one of the allocated frames, and the frame number is put in the **page table** for this process.
- The **next page** is loaded into another frame, its **frame number** is recorded in the page table, and so on (**Figure 9.11**).



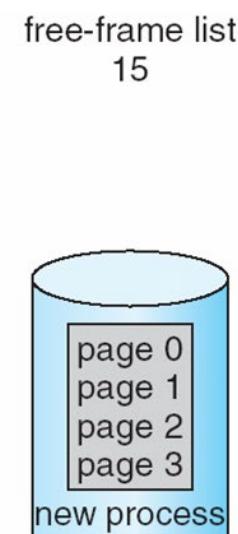
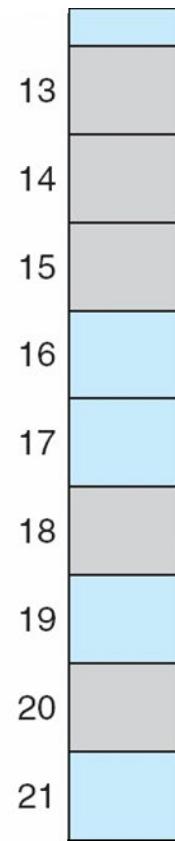


# Paging



(a)

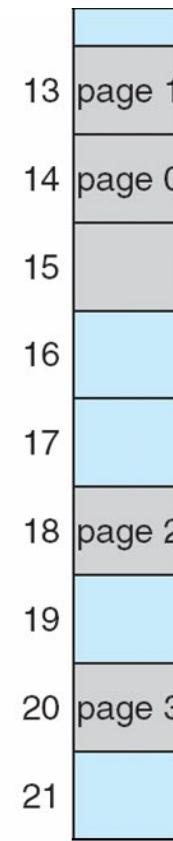
Before allocation



new-process page table

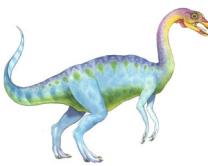
(b)

After allocation



**Figure 9.11** Free frames (a) before allocation and (b) after allocation.





# Paging

---

- An important aspect of paging is the clear separation between the **programmer's view of memory** and the **actual physical memory**.
  - *The programmer views memory as one single space, containing only this one program.*
- In fact, the user program is scattered throughout physical memory, which also holds other programs.
  - The difference between the **programmer's view of memory** and the **actual physical memory** is reconciled by the **address-translation hardware**.
  - The logical addresses are translated into physical addresses.
  - This mapping is hidden from the programmer and is controlled by the OS.





# Paging

- Since the **OS** is managing physical memory:
  - *The OS determines which frames are allocated, which frames are available, how many total frames there are, and so on.*
  - This information is generally kept in a single, system-wide data structure of the OS, which is called a **frame table**.
  - If a user makes a **system call** (to do I/O, for example) and provides an address as a parameter (a buffer, for instance), that address must be mapped to produce the correct **physical address**.
- The OS maintains a **copy of the page table (PT)** for each process:
  - This **PT** copy is used **to translate logical addresses to physical addresses** whenever the OS must map a logical address to a physical address manually.
  - The CPU dispatcher defines a **hardware page table** when a process is allocated the CPU. Therefore, Paging increases the context-switch time.





# Implementation of Page Table

---

- **Page table** is kept in **main memory**.
- **Page-table base register (PTBR)** points to the **page table**.
- **Page-table length register (PTLR)** indicates **size of the page table**.
- In this scheme every data/instruction access **requires two memory accesses**
  - One for the **page table** and one for the **data / instruction from physical memory**
- The **two-memory access problem** can be solved using a **special fast-lookup hardware cache** called **translation look-aside buffers (TLBs)**.





# Implementation of Page Table

- If the **page number** is not in the **TLB** (known as a **TLB miss**), a **memory** reference to the **page table** must be made.
- Depending on the CPU, this may be done automatically in hardware or via an **interrupt to the operating system**.
- When the **frame number** is obtained, we can use it to access memory (**Figure 8.14**).
- In addition, we add the **page number** and **frame number** to the **TLB**, so that they will be found quickly on the next reference. If the TLB is already full of entries, an existing entry must be selected for **replacement**.
- **Replacement** policies range from **least recently used (LRU)** through round-robin to random.
- Some CPUs allow the operating system to participate in LRU entry replacement, while others handle the matter themselves.





# Implementation of Page Table

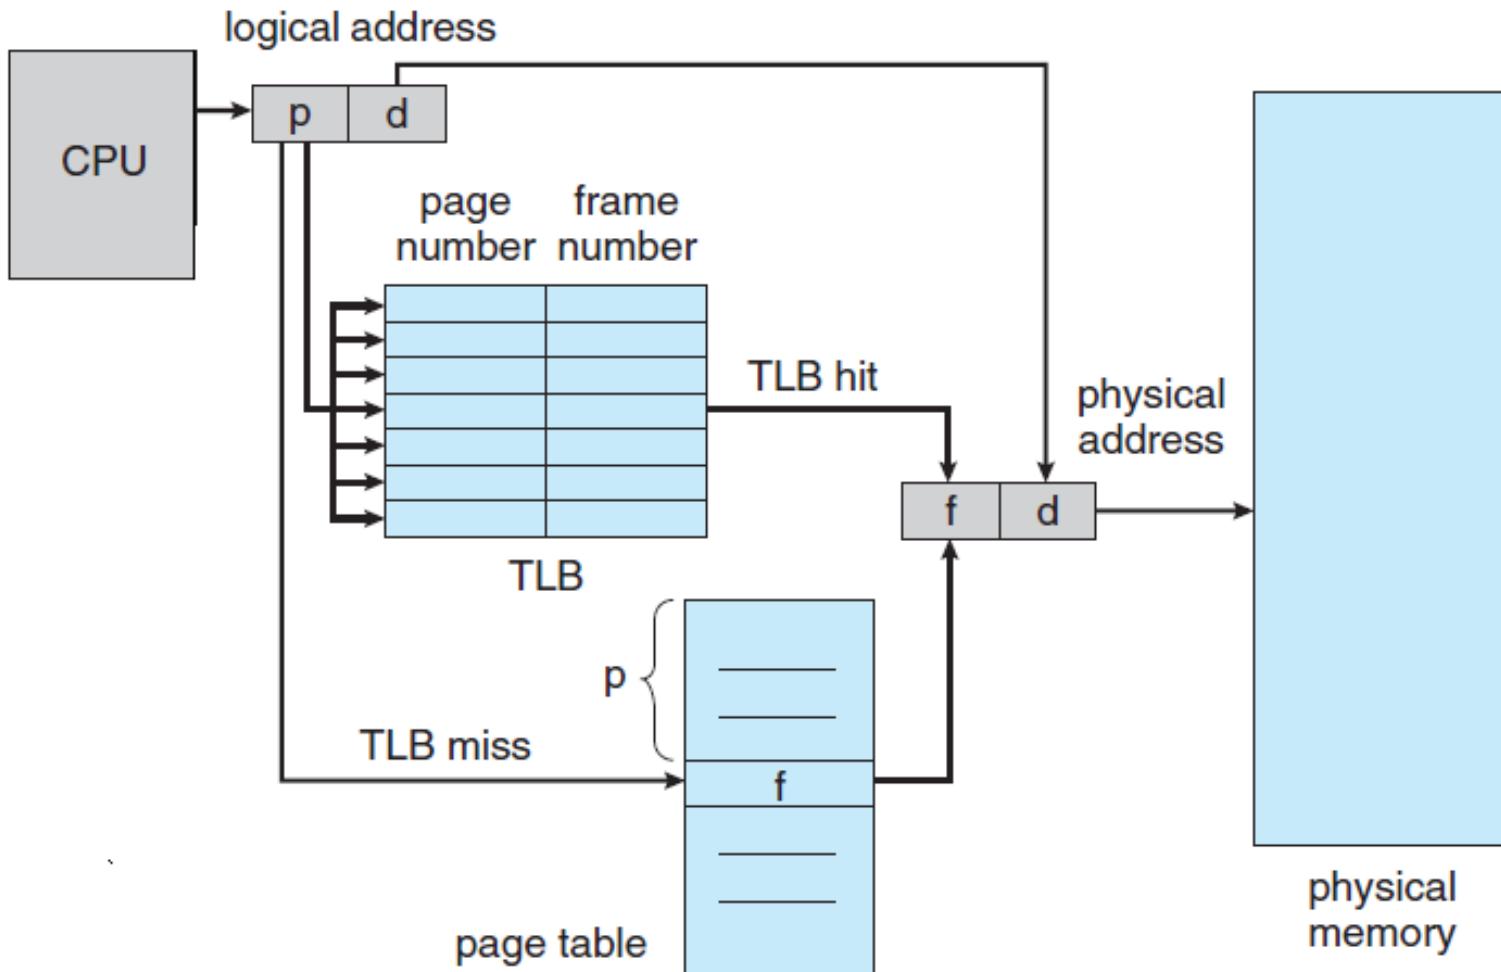
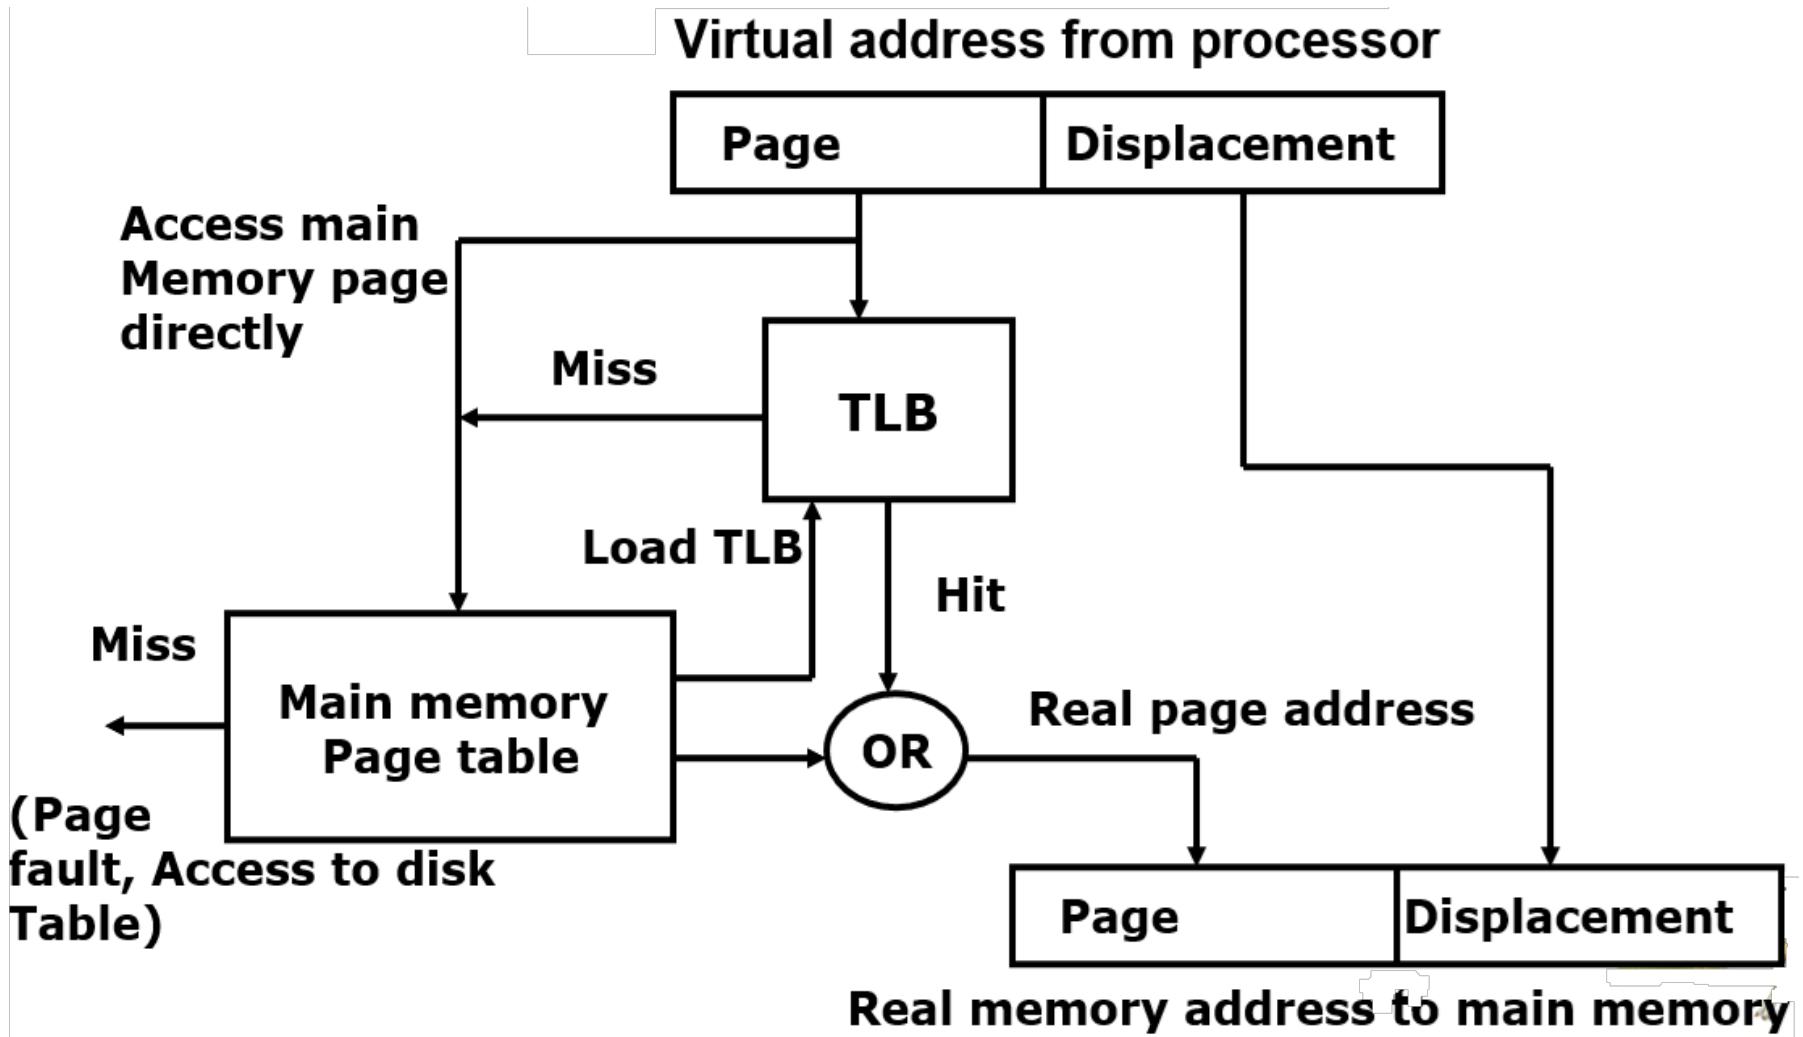


Figure 8.14 Paging hardware with TLB.





# Implementation of Page Table





# Implementation of Page Table

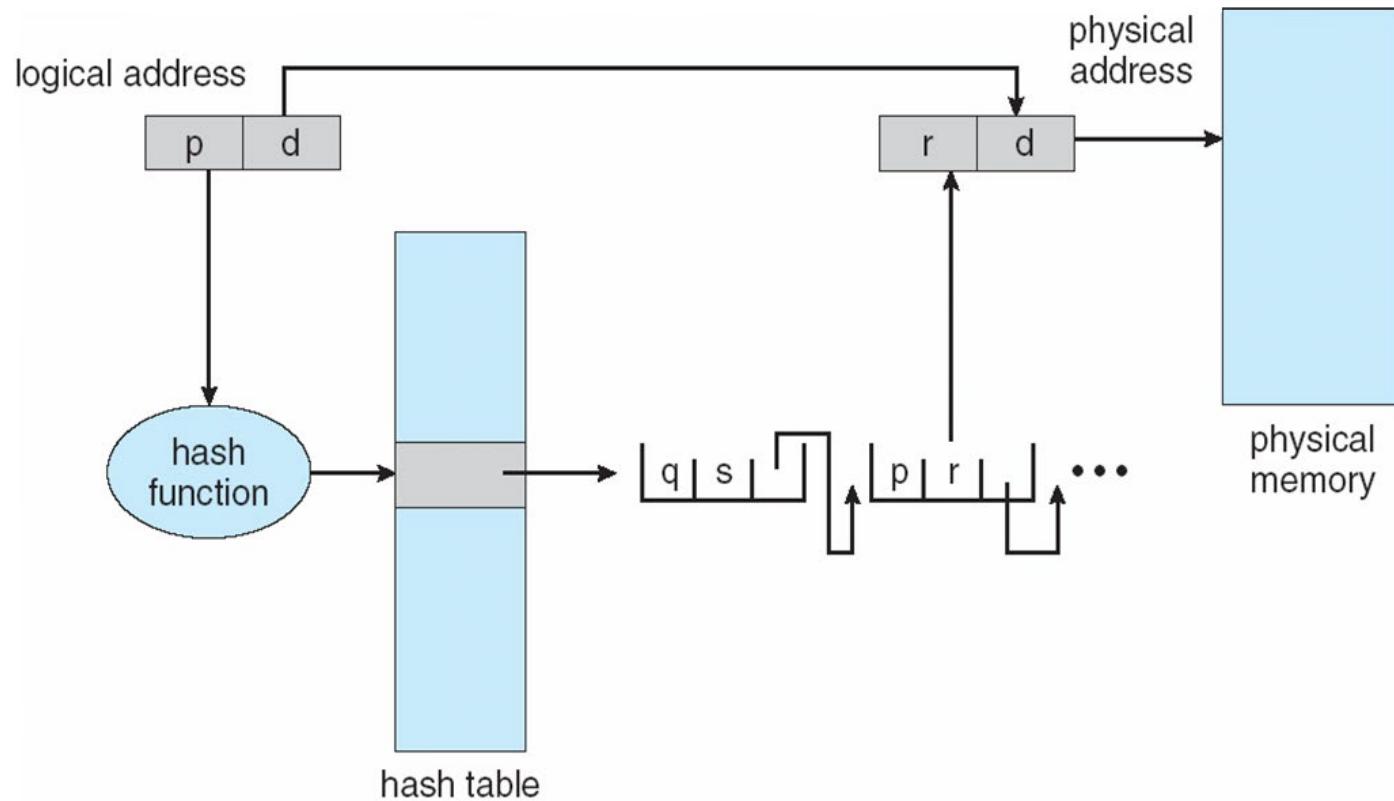
---

- The percentage of times that the **page number** of interest is **found in the TLB** is called the **hit ratio**.
- An **80-percent hit ratio**, for example, means that the desired **page number** stays in the TLB 80 percent of the time.
- If it takes **10 nanoseconds to access memory**, then a mapped-memory access takes **10 nanoseconds** when the **page number** is in the TLB.
- If it fails to find the **page number** in the TLB, then **first access memory for the page table and frame number** (10 nanoseconds) and then **next access the desired byte in memory** (10 nanoseconds), for a total of **20 nanoseconds**.
- **Effective memory-access time = (hit-percent × memory\_access\_time) + (miss-percent × page\_miss\_handling\_time)**
- **The effective access time =  $(0.80 \times 10) + (0.20 \times 20) = 12 \text{ ns}$**
- **Find an effective memory access time when the hit ratio is 99%**





# Hashed Page Table





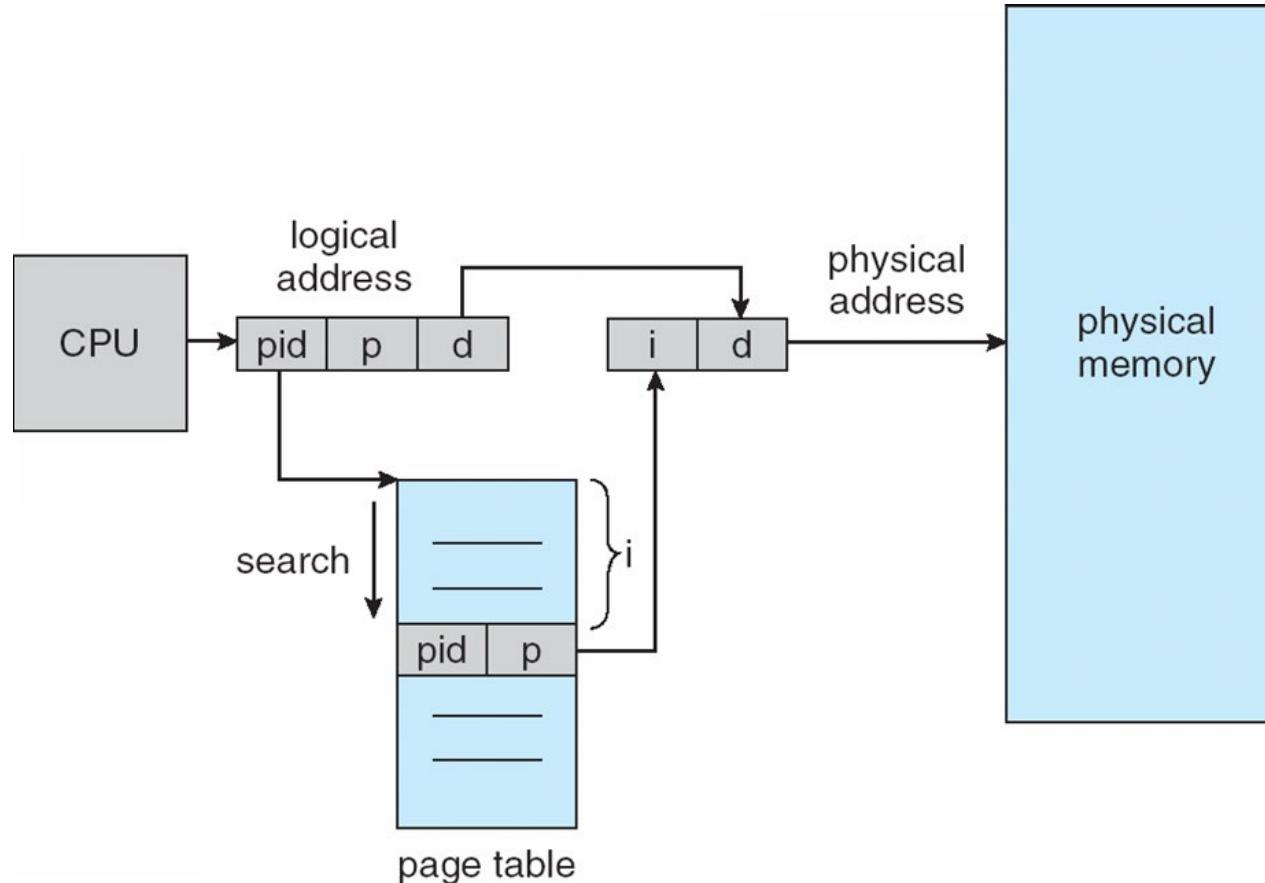
# Inverted Page Table

- Rather than each process having a **page table** and keeping track of all possible **logical pages**, track all **physical pages**
- An **inverted page table** has one table entry for every page frame in real memory, resulting significant **reduction in page table size**
- One entry for each real page of memory
- Entry consists of the **virtual address** of the page stored in that **real memory location**, with information about the process that owns that page
- Decreases memory needed to store each **page table**, but increases time needed to search the table when a page reference occurs





# Inverted Page Table Architecture





# Inverted Page Table Architecture

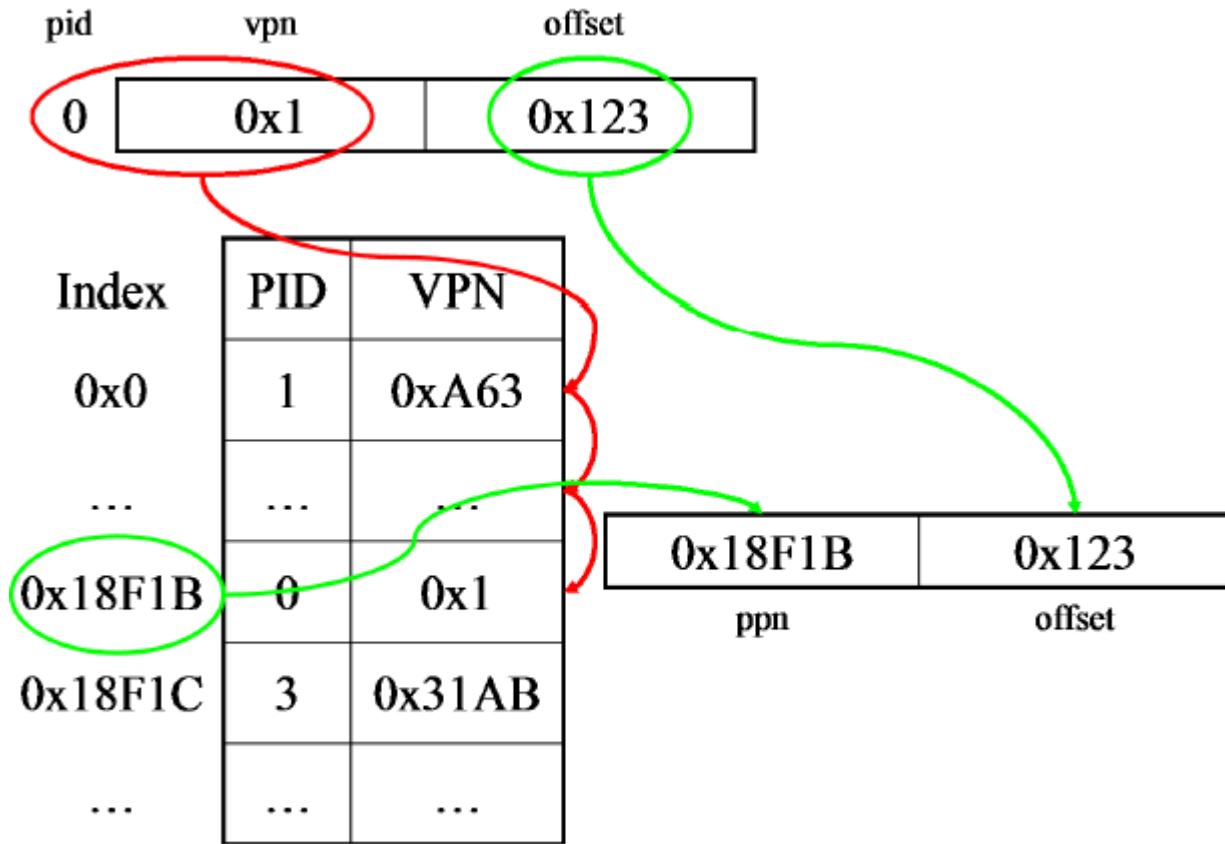


Figure : Translation procedure using a linear inverted page table.  
Info bits exist in each entry, though they are not shown.





# Page Protection

- **Memory protection** in a paged environment is accomplished by protection bits associated with each frame.
- Normally, these bits are kept in the **page table**.
- **One bit** can define a page to be **read-write** or **read-only**.
- Every reference to memory goes through the **page table** to find the **correct frame number**.
- At the same time, the **physical address** is being computed, the **protection bits** can be checked to verify that no writes are being made to a read-only page.
- An attempt to write to a **read-only page** causes a **hardware trap** to the operating system (or memory-protection violation).

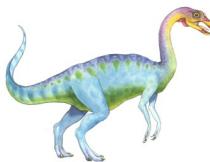




# Page Protection

- One **additional bit** is generally attached to each entry in the **page table**: a **valid –invalid** bit.
- When this bit is set to **valid**, the associated page is in the process's logical address space and is thus a legal (or valid) page.
- When the bit is set to **invalid**, the page is not in the process's logical address space.
- Illegal addresses are trapped using the **valid–invalid bit**.
- The OS sets this bit for each page to allow or disallow access to the page.





# Page Protection

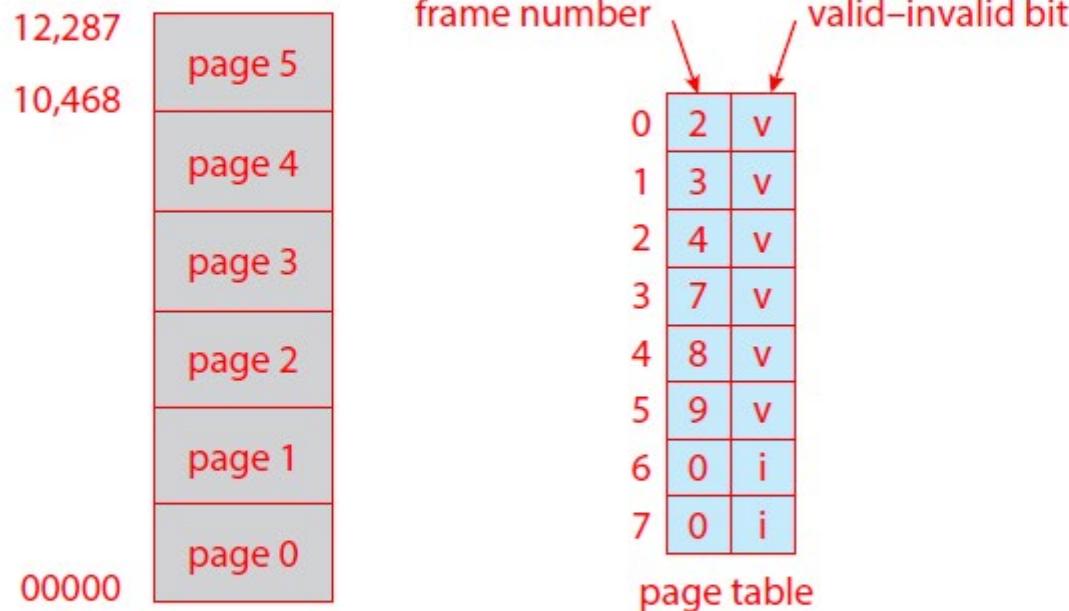
---

- Suppose, for example, that in a system with a **14-bit address space** (0 to 16383), we have a program that should use only addresses **0 to 10468**.
- Given a page size of **2 KB**, we have the situation shown in **Figure 9.13**.
- Addresses in **pages 0, 1, 2, 3, 4, and 5** are mapped normally through the **page table**.
- Any attempt to generate an address on **pages 6 or 7**, however, will find that the **valid–invalid bit** is set to **invalid**, and the computer will **trap** to the OS (it is an **invalid page reference**).

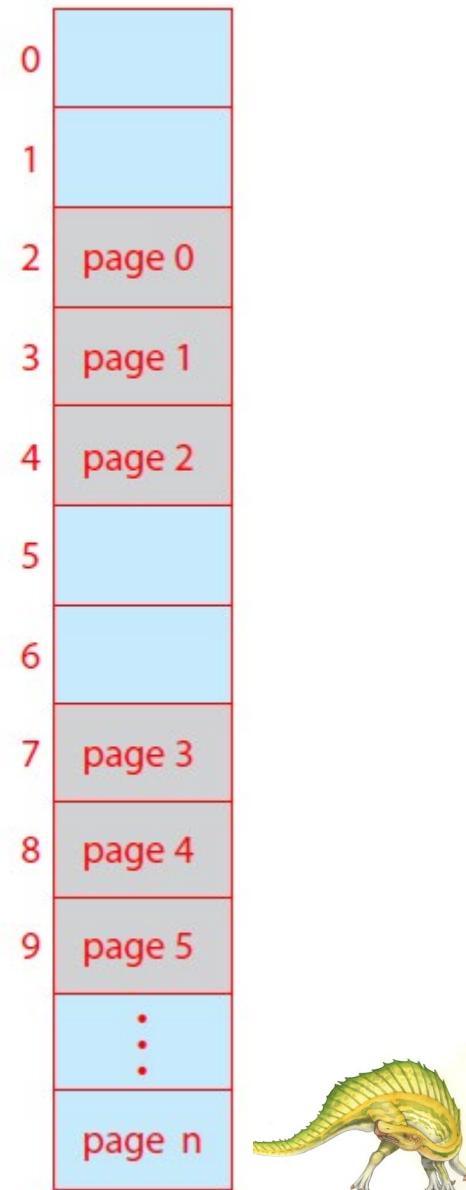




# Page Protection



**Figure 9.13** Valid (v) or invalid (i) bit in a page table.





# Page Protection

- Notice that this scheme has created a problem.
- Because the program extends only to address **10468**, any reference beyond that address is illegal.
- However, **references to page 5** are classified as **valid**, so accesses to **addresses up to 12287** are **valid**.
- Only the **addresses from 12288 to 16383** are **invalid**.
- This problem is due to the **2-KB page size** and reflects **internal fragmentation of paging**.





# Swapping

---

- *Instructions and the data of a process must be in the main memory to be executed.*
- However, a process, or a portion of a process, can be **swapped** temporarily out of memory to a **backing store** and then brought back into memory for continued execution (**Figure 9.19**).
- **Swapping** allows the total **physical address space** of all processes to exceed the real physical memory of the system, thus *increasing the degree of multiprogramming in a system.*





# Swapping

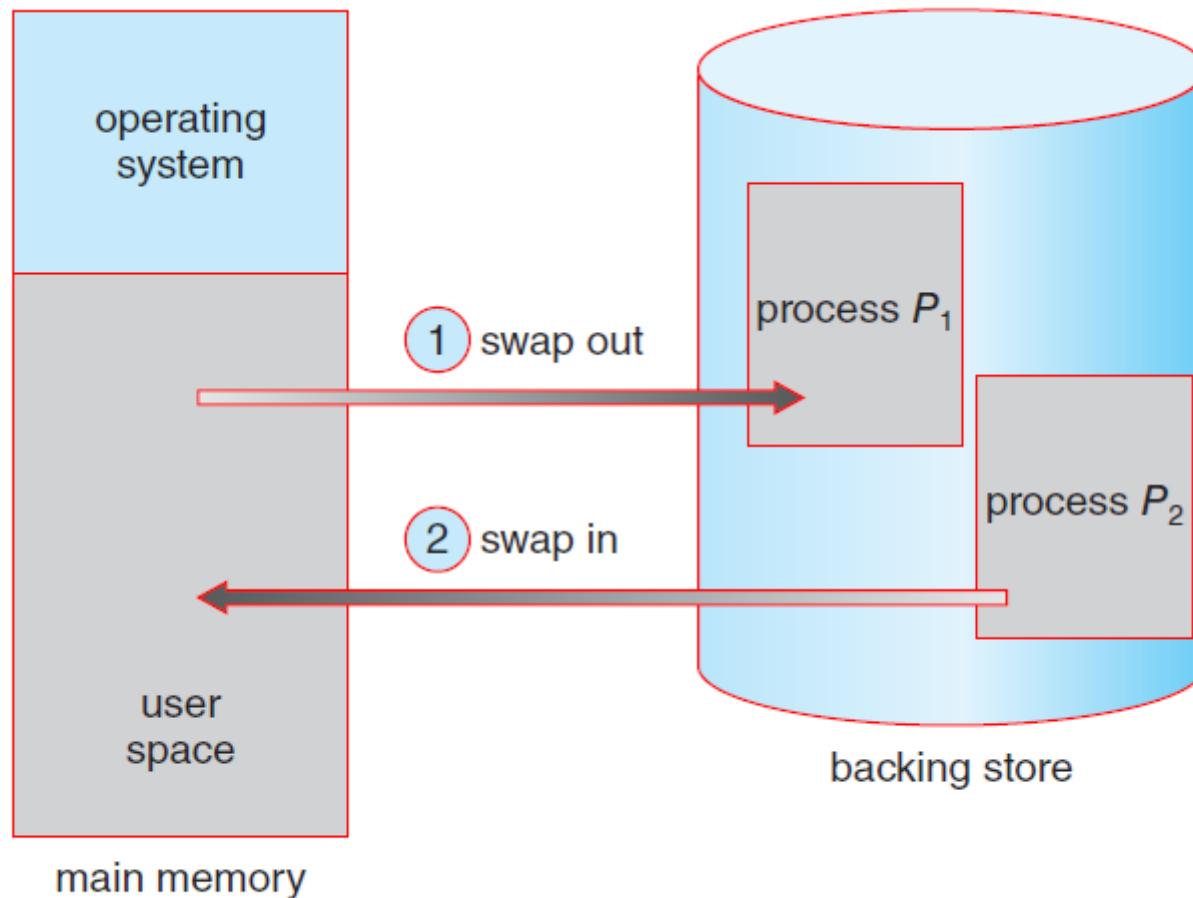


Figure 9.19 Standard swapping of two processes using a disk as a backing store.





# Swapping

---

- Standard swapping was used in traditional UNIX systems.
  - But it is generally no longer used in contemporary operating systems, because the amount of time required to move entire processes between memory and the backing store is prohibitive.
- Most operating systems, including **Linux** and **Windows**, now use a variation of swapping in which **pages** of a process, rather than the entire process, are swapped.
- In **swapping with paging**, the **page-out** operation moves a page from memory to the backing store; the reverse process is known as a **page-in**.
- Swapping with paging is illustrated in **Figure 9.20**, where a subset of pages for **processes A and B** is being **paged out** and **paged in**, respectively.





# Swapping

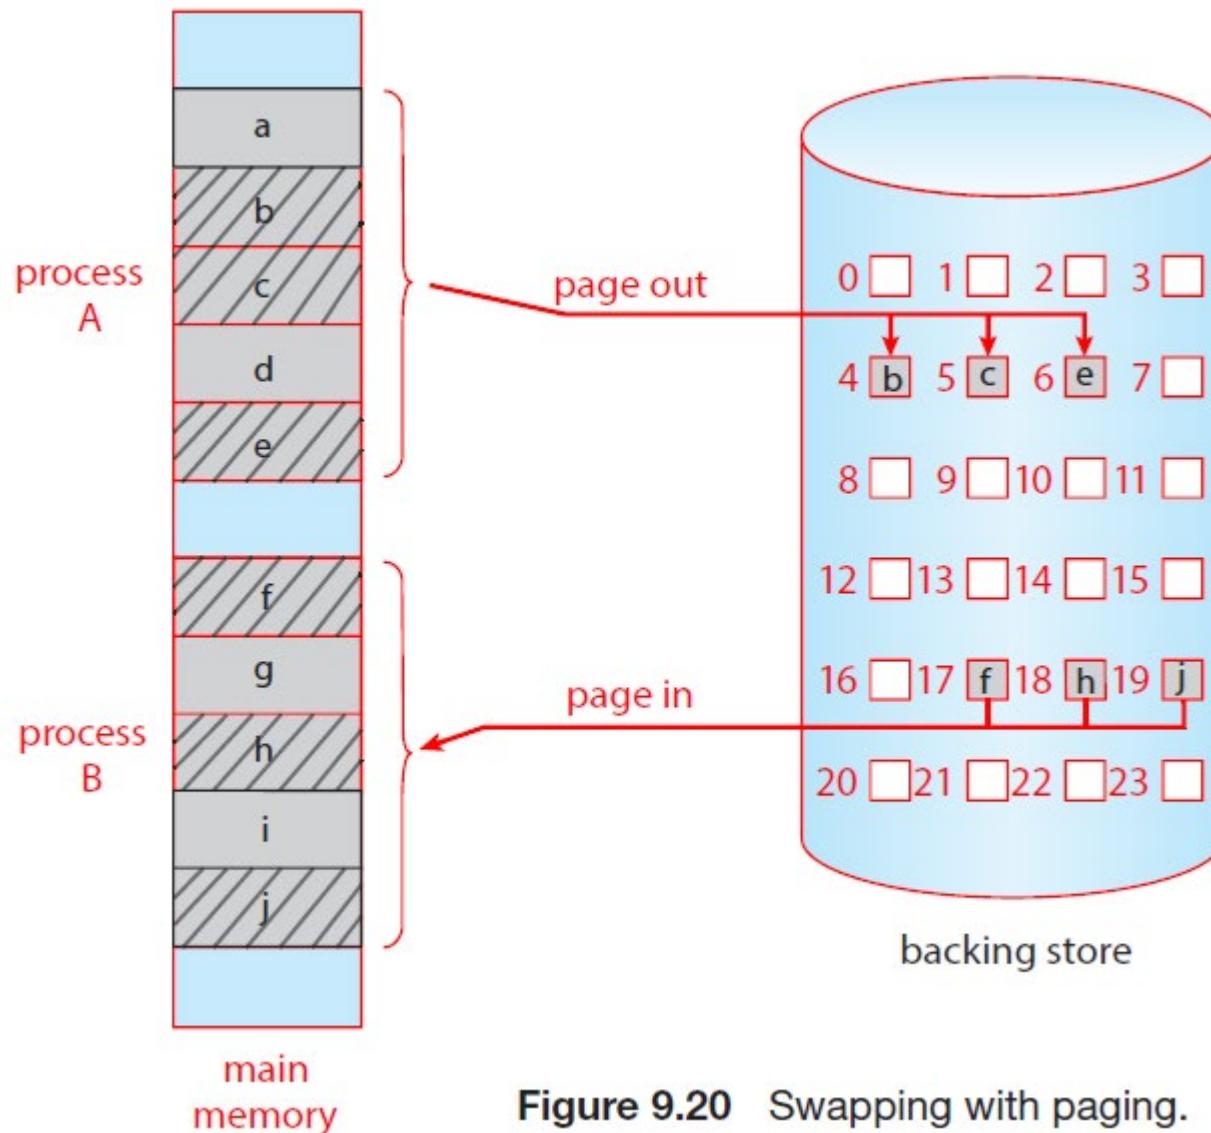
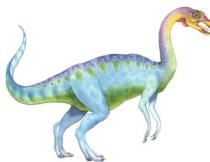


Figure 9.20 Swapping with paging.





# Example: ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed **sections**)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
  - Outer level has two micro TLBs (one data, one instruction)
  - Inner is single main TLB
  - First inner is checked, on missouters are checked, and on miss page table walk performed by CPU

